

Predicting Vehicle Prices Using Regression Models

Objectives

The primary goal of this project is to develop a robust regression model to predict used car prices for a reseller based on various listed features and specifications. In addition to predicting prices, the project focuses on identifying feature importance and mitigating overfitting through the application of regularisation techniques.

There can be several business objectives for this, such as:

- **Price Prediction:** Model car prices based on features like mileage, fuel type, and performance.
- **Market Analysis:** Explore trends and preferences in the used car market, by type, region, or other metrics.
- **Feature Importance:** Identify the most important factors influencing car prices (e.g., fuel type, mileage, age).

Tasks Overview

The data pipeline for this task involves the following steps:

1. **Dataset Overview**
2. **Data Preprocessing**
3. **Data Visualisation & Exploration**
4. **Model Building**
5. **Regularisation**

1 Data Understanding

| **Variable | Description** | -----|-----| | `make_model` | The brand and model of the vehicle (e.g., 'Audi A1'). | | `body_type` | The body style of the vehicle, such as Sedan, Compact, or Station Wagon. | | `price` | The listed price of the car in currency. | | `vat` | Indicates the VAT status for the vehicle's price (e.g., VAT deductible, Price negotiable). | | `km` | The total mileage (in kilometers) of the vehicle, indicating its usage. | | `Type` | Condition of the vehicle, whether it's 'Used' or 'New'.| | `Fuel` | Type of fuel the vehicle uses, such as 'Diesel', 'Benzine', etc. | | `Gears` | The number of gears in the vehicle's transmission. | | `Comfort_Convenience` | Comfort and convenience features, such as 'Air conditioning', 'Leather steering wheel', 'Cruise control', and more. | |

`Entertainment_Media` | Media features available in the vehicle, including 'Bluetooth', 'MP3', 'Radio', etc. || `Extras` | Additional features like 'Alloy wheels', 'Sport suspension', etc.|| `Safety_Security` | Safety features like 'ABS', 'Airbags', 'Electronic stability control', 'Isofix', etc. || `age` | Age of the car (calculated based on the model year). || `Previous_Owners` | The number of previous owners the car has had. || `hp_kw` | Engine power in kilowatts (kW), indicating the performance capacity of the engine.|| `Inspection_new` | Indicates whether the car has recently undergone an inspection (1 for yes, 0 for no). || `Paint_Type` | The type of paint on the car, such as 'Metallic', 'Matte', etc. || `Upholstery_type` | The material used for the interior upholstery, such as 'Cloth', 'Leather', etc.|| `Gearing_Type` | The type of transmission the car uses, either 'Automatic' or 'Manual'. || `Displacement_cc` | The engine displacement in cubic centimeters (cc), indicating the size of the engine.|| `Weight_kg` | The total weight of the vehicle in kilograms.|| `Drive_chain` | The type of drivetrain, indicating whether it's 'Front' or 'Rear' wheel drive. || `cons_comb` | The combined fuel consumption in liters per 100 kilometers.|

1.1 Data Loading

Importing Necessary Libraries

```
In [121...]: # Importing necessary libraries
# Core Python
import numpy as np
import pandas as pd

# Visualization
import matplotlib.pyplot as plt

# Model building & preprocessing
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline

# Regression models
from sklearn.linear_model import LinearRegression, RidgeCV, Lasso, ElasticNet

# Model evaluation
from sklearn.metrics import mean_squared_error, r2_score
```

1.1.1

Load the dataset

```
In [122...]: # Load the data
# Load the dataset
import pandas as pd

file_path = "Car_Price_data.csv" # Update path if your file is in another c
```

```

df = pd.read_csv(file_path)

print("Data loaded successfully.")
print("Shape:", df.shape)

# Preview first few rows
df.head()

```

Data loaded successfully.
Shape: (15915, 23)

Out[122...]

	make_model	body_type	price	vat	km	Type	Fuel	Gears
0	Audi A1	Sedans	15770	VAT deductible	56013.0	Used	Diesel	7.0
1	Audi A1	Sedans	14500	Price negotiable	80000.0	Used	Benzine	7.0
2	Audi A1	Sedans	14640	VAT deductible	83450.0	Used	Diesel	7.0
3	Audi A1	Sedans	14500	VAT deductible	73000.0	Used	Diesel	6.0
4	Audi A1	Sedans	16790	VAT deductible	16200.0	Used	Diesel	7.0

5 rows × 23 columns

2 Analysis and Feature Engineering [35 marks]

2.1 Preliminary Analysis and Frequency Distributions [13 marks]

2.1.1 [1 marks]

Check and fix missing values.

In [123...]

```

# Find the proportion of missing values in each column and handle if found

missing_pct = df.isnull().mean().sort_values(ascending=False)

print("Missing value percentage per column:")
print(missing_pct)

numeric_cols = df.select_dtypes(include=['int64', 'float64']).columns
categorical_cols = df.select_dtypes(include=['object', 'category']).columns
df[numeric_cols] = df[numeric_cols].fillna(df[numeric_cols].median())

```

```
df[categorical_cols] = df[categorical_cols].fillna(df[categorical_cols].mode)

print("\nMissing values handled successfully.")

Missing value percentage per column:
make_model          0.0
age                 0.0
Drive_chain         0.0
Weight_kg           0.0
Displacement_cc     0.0
Gearing_Type        0.0
Upholstery_type    0.0
Paint_Type          0.0
Inspection_new     0.0
hp_kw               0.0
Previous_Owners    0.0
Safety_Security     0.0
body_type           0.0
Extras              0.0
Entertainment_Media 0.0
Comfort_Convenience 0.0
Gears               0.0
Fuel                0.0
Type                0.0
km                  0.0
vat                 0.0
price               0.0
cons_comb           0.0
dtype: float64
```

Missing values handled successfully.

From the features, identify the target feature and numerical and categorical predictors. Select the numerical and categorical features carefully as they will be used in analysis.

2.1.2 [3 marks]

Identify numerical predictors and plot their frequency distributions.

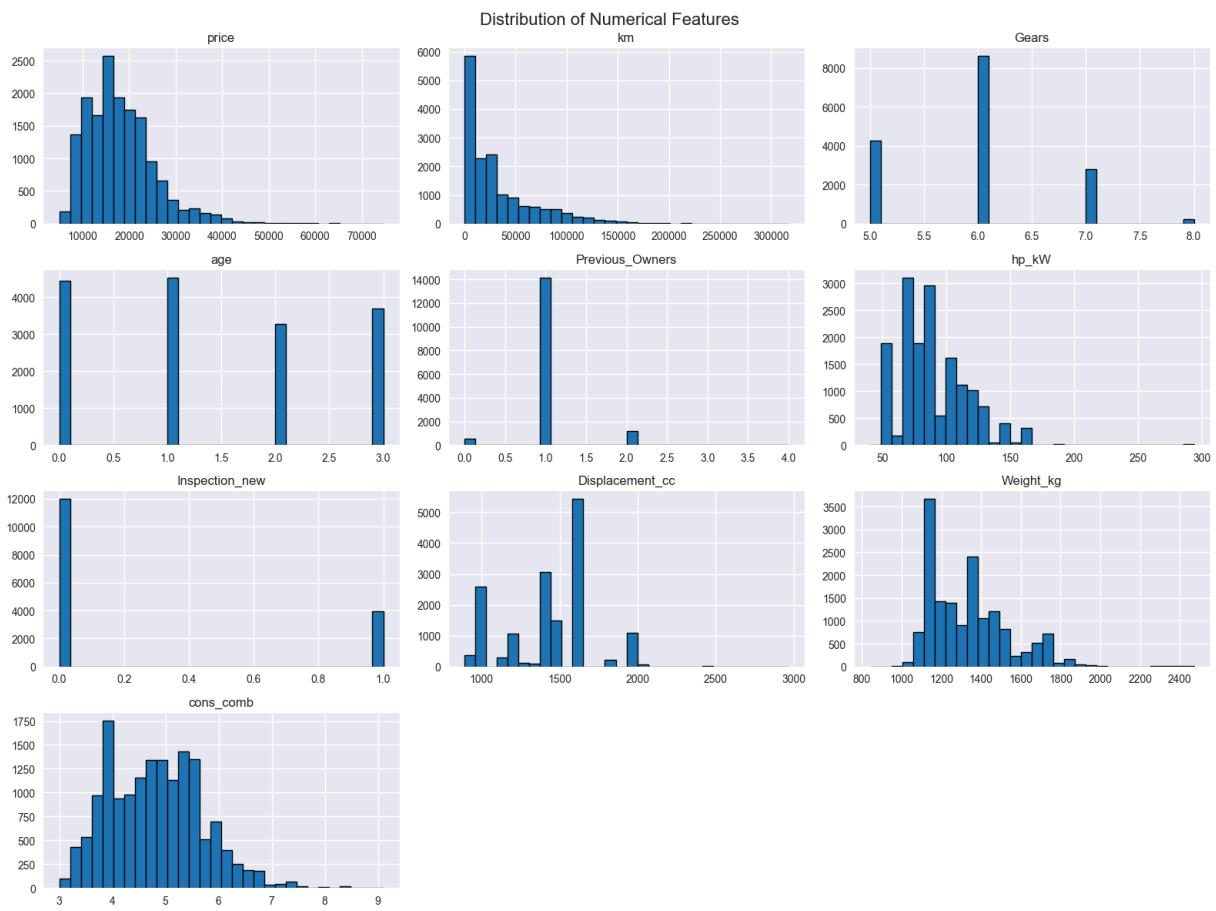
```
In [124...]: # Identify numerical features and plot histograms

import matplotlib.pyplot as plt

numeric_cols = df.select_dtypes(include=['int64', 'float64']).columns
print("Numerical features:", list(numeric_cols))

df[numeric_cols].hist(figsize=(16, 12), bins=30, edgecolor='black')
plt.suptitle("Distribution of Numerical Features", fontsize=16)
plt.tight_layout()
plt.show()
```

Numerical features: ['price', 'km', 'Gears', 'age', 'Previous_Owners', 'hp_kw', 'Inspection_new', 'Displacement_cc', 'Weight_kg', 'cons_comb']



2.1.3 [3 marks]

Identify categorical predictors and plot their frequency distributions.

In [125...]

```
# Identify categorical columns and check their frequency distributions

# Identify categorical predictors
categorical_cols = df.select_dtypes(include=['object', 'category']).columns.

# Columns that should NOT be treated as normal categorical features
multi_label_cols = [
    "Comfort_Convenience",
    "Entertainment_Media",
    "Extras",
    "Safety_Security"
]

# Final categorical predictors
categorical_predictors = [col for col in categorical_cols if col]

print("Categorical predictors:", categorical_predictors)

import matplotlib.pyplot as plt

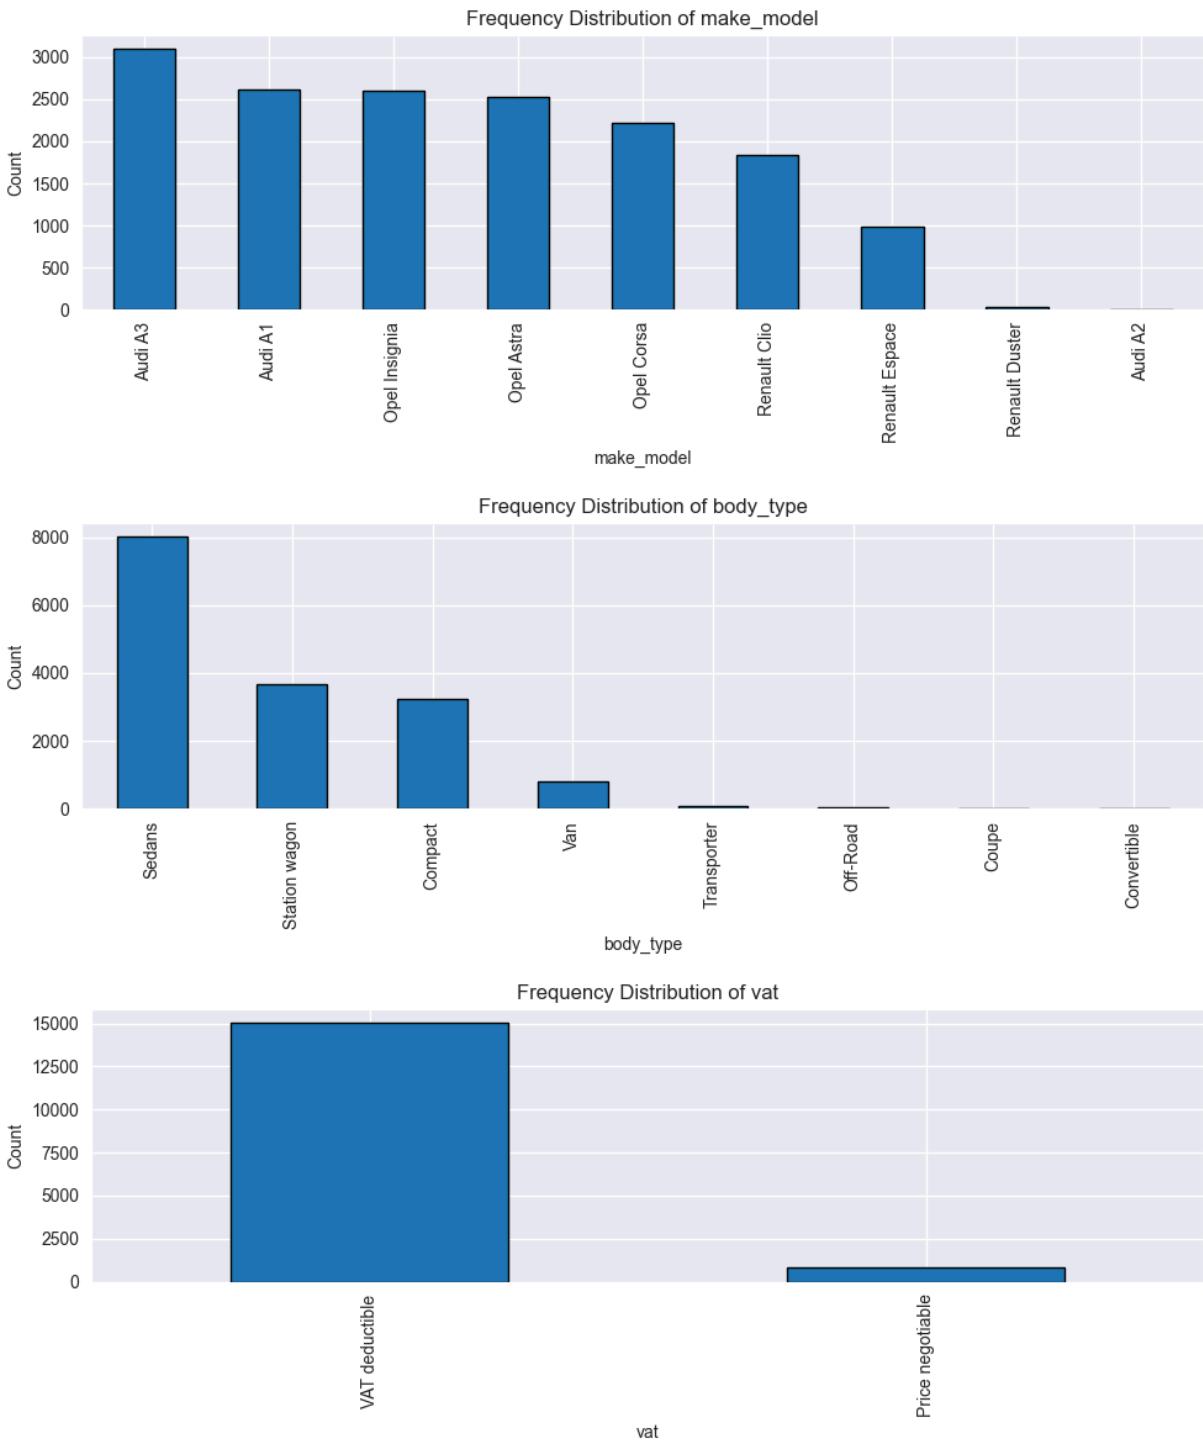
for col in categorical_predictors:
    plt.figure(figsize=(10, 4))
```

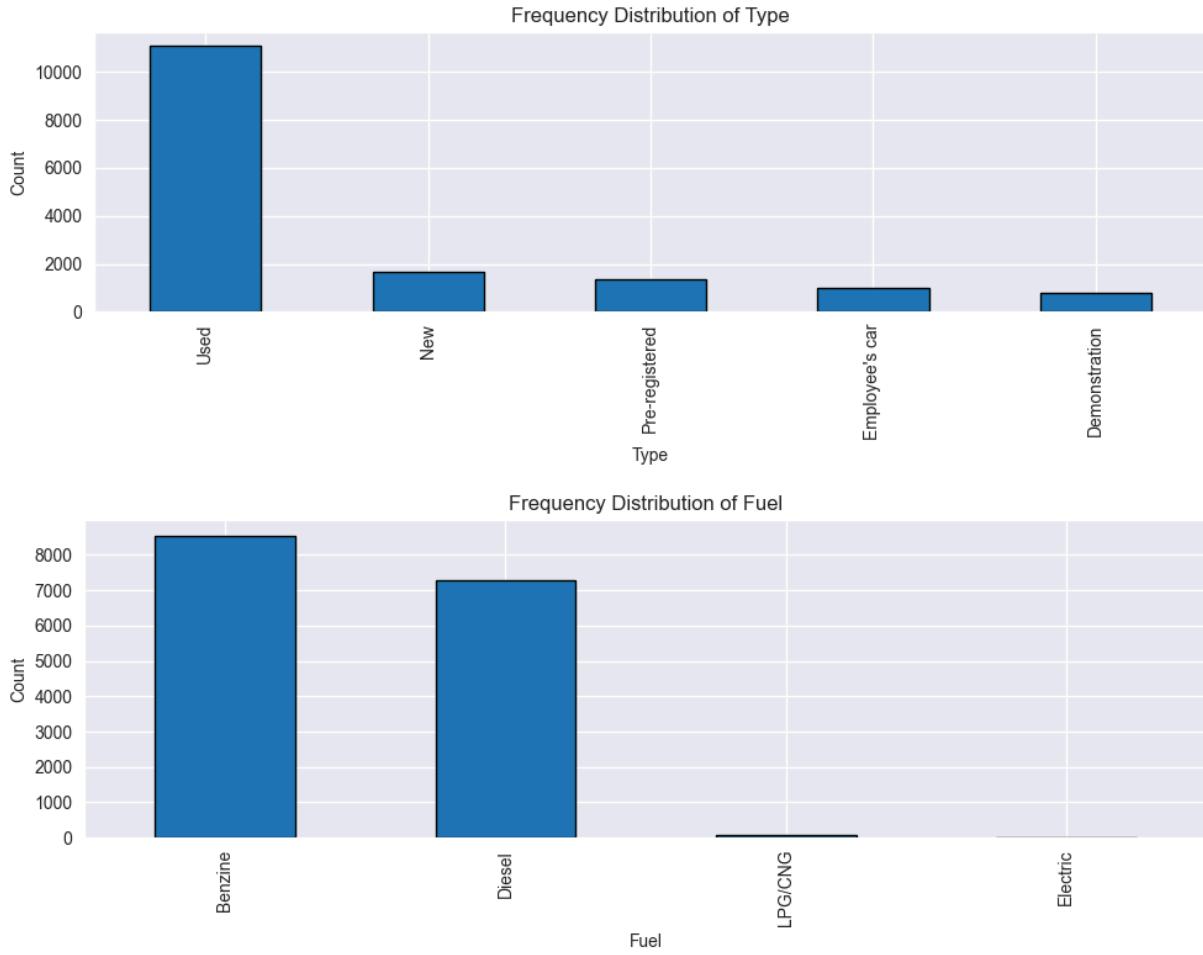
```

df[col].value_counts().head(20).plot(kind='bar', edgecolor='black')
plt.title(f"Frequency Distribution of {col}")
plt.xlabel(col)
plt.ylabel("Count")
plt.tight_layout()
plt.show()

```

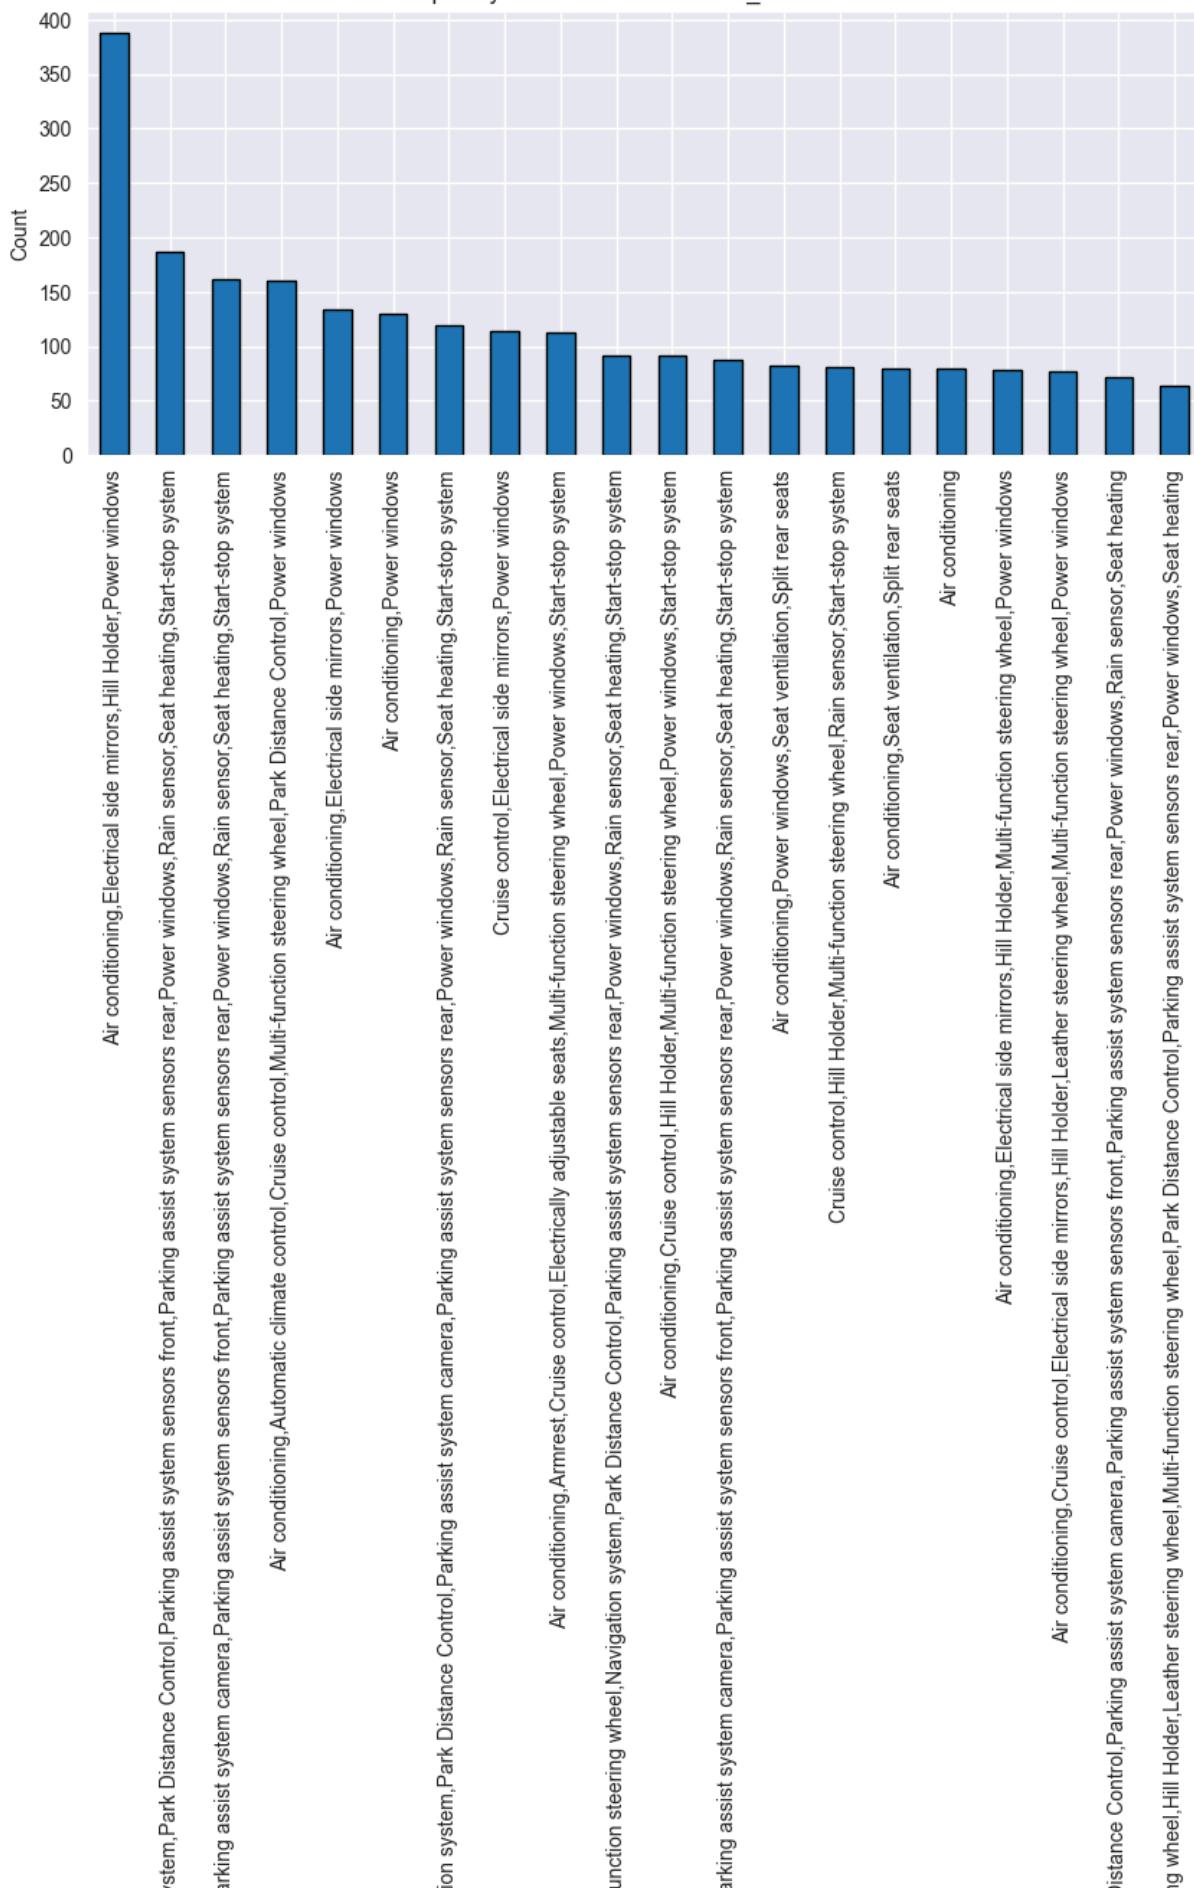
Categorical predictors: ['make_model', 'body_type', 'vat', 'Type', 'Fuel', 'Comfort_Convenience', 'Entertainment_Media', 'Extras', 'Safety_Security', 'Paint_Type', 'Upholstery_type', 'Gearing_Type', 'Drive_chain']





```
/var/folders/q9/_3w68hr55rdcl70s4bt4j_jbh0000gn/T/ipykernel_66612/2463149345.py:28: UserWarning: Tight layout not applied. The bottom and top margins can not be made large enough to accommodate all Axes decorations.  
plt.tight_layout()
```

Frequency Distribution of Comfort_Convenience



Air conditioning,Cruise control,Electrically heated windshield,Electrical side mirrors,Heated steering wheel,Hill Holder,Leather steering wheel,Light sensor,Multi-function steering wheel,Navigation system,Park Distance Control,Park D

Air conditioning,Cruise control,Electrical side mirrors,Heated side mirrors,Heated steering wheel,Light sensor,Multi-function steering wheel,Navigation system,Park Distance Control,Pc

Air conditioning,Armrest,Automatic climate control,Cruise control,Electric side mirrors,Heated steering wheel,Hill Holder,Keyless central door lock,Leather steering wheel,Light sensor,Lumbar support,Multi-function steering wheel,Navigation system,Park Distance Control,Pc

Air conditioning,Armrest,Automatic climate control,Cruise control,Electric side mirrors,Heated steering wheel,Hill Holder,Leather steering wheel,Light sensor,Multi-function steering wheel,Navigation system,Park Distance Control,Pc

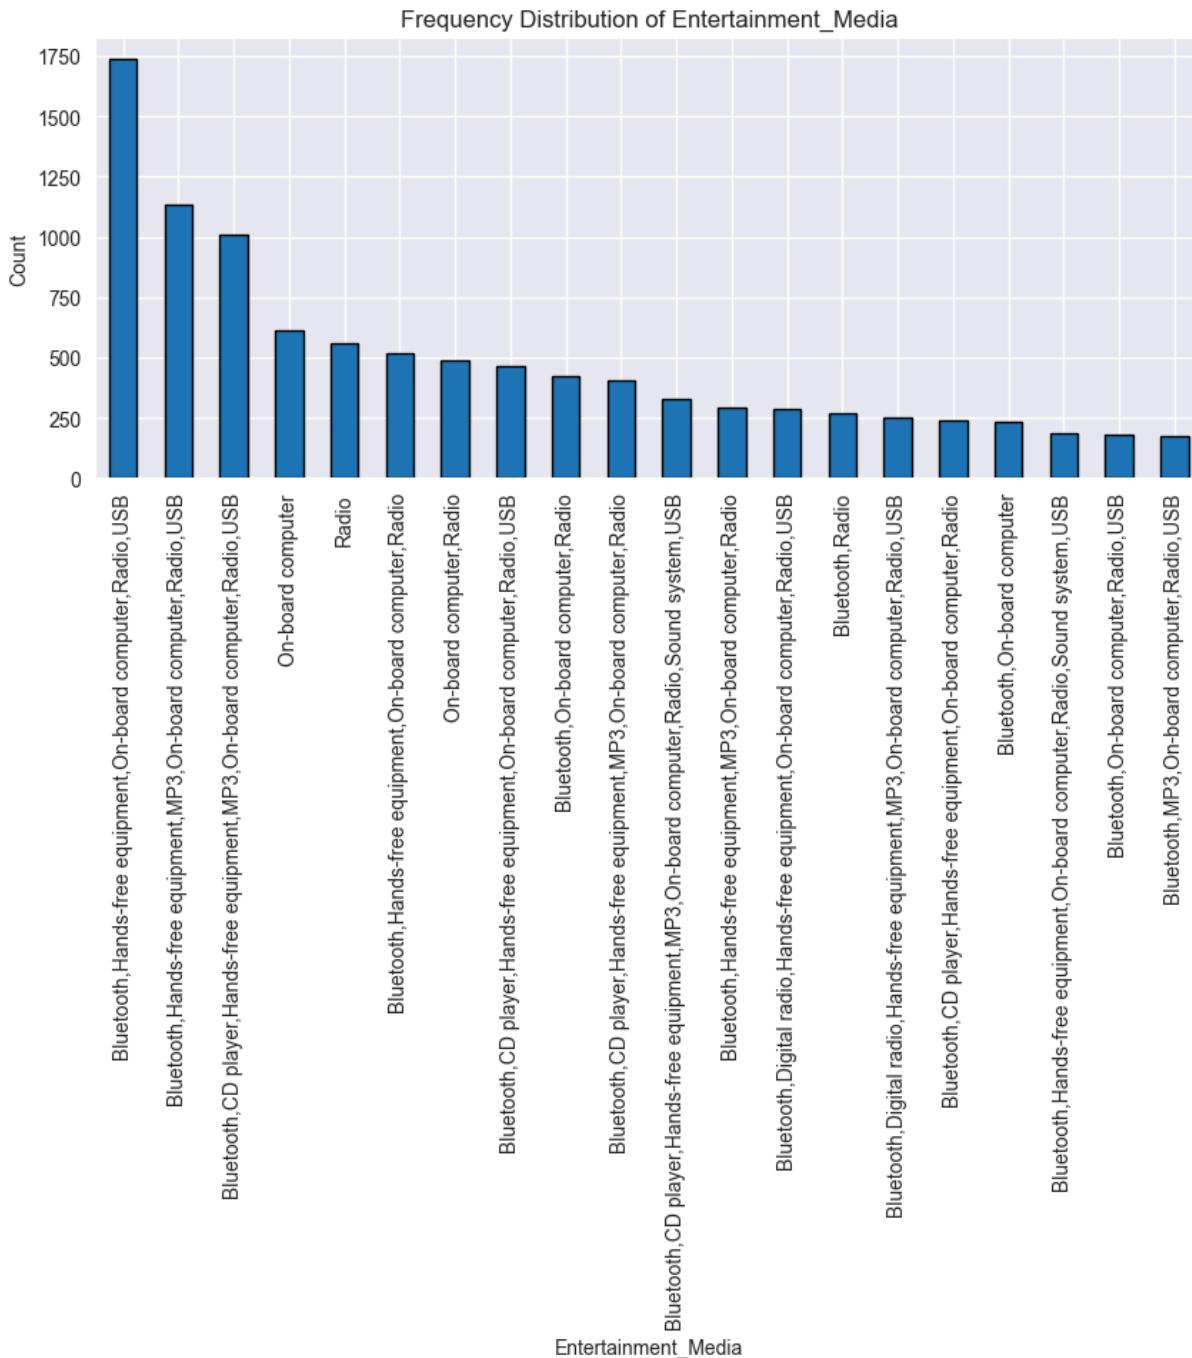
Air conditioning,Armrest,Automatic climate control,Cruise control,Electric side mirrors,Heated steering wheel,Hill Holder,Leather steering wheel,Light sensor,Multifunction steering wheel,Navigation system,Park Distance Control,Pc

Air conditioning,Armrest,Automatic climate control,Cruise control,Electrically adjustable seats,Electric

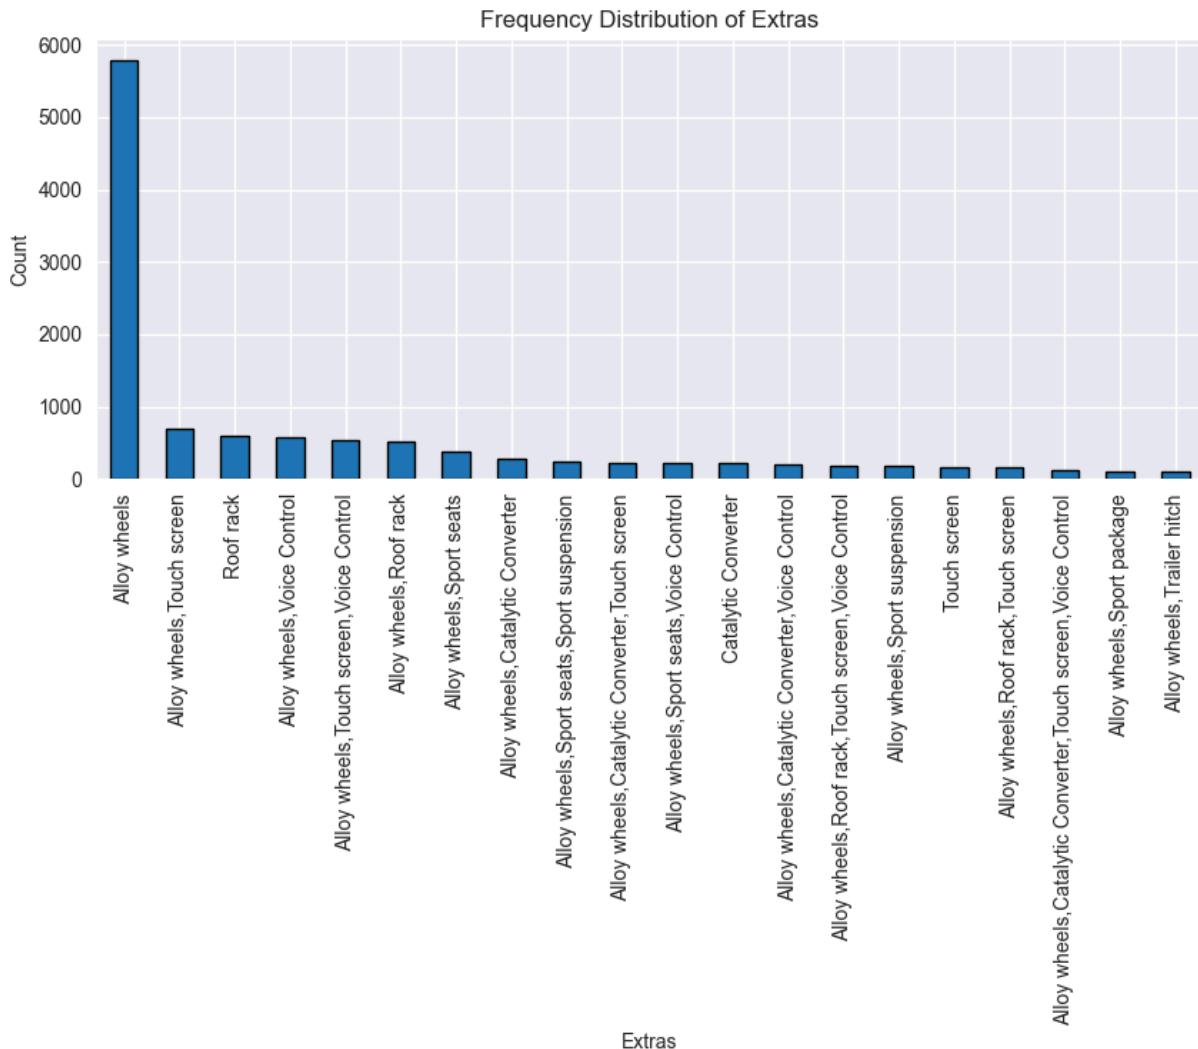
Air conditioning,Armrest,Automatic climate contr

Comfort_Convenience

```
/var/folders/q9/_3w68hr55rdcl70s4bt4jjbh0000gn/T/ipykernel_66612/2463149345.py:28: UserWarning: Tight layout not applied. The bottom and top margins can not be made large enough to accommodate all Axes decorations.  
plt.tight_layout()
```

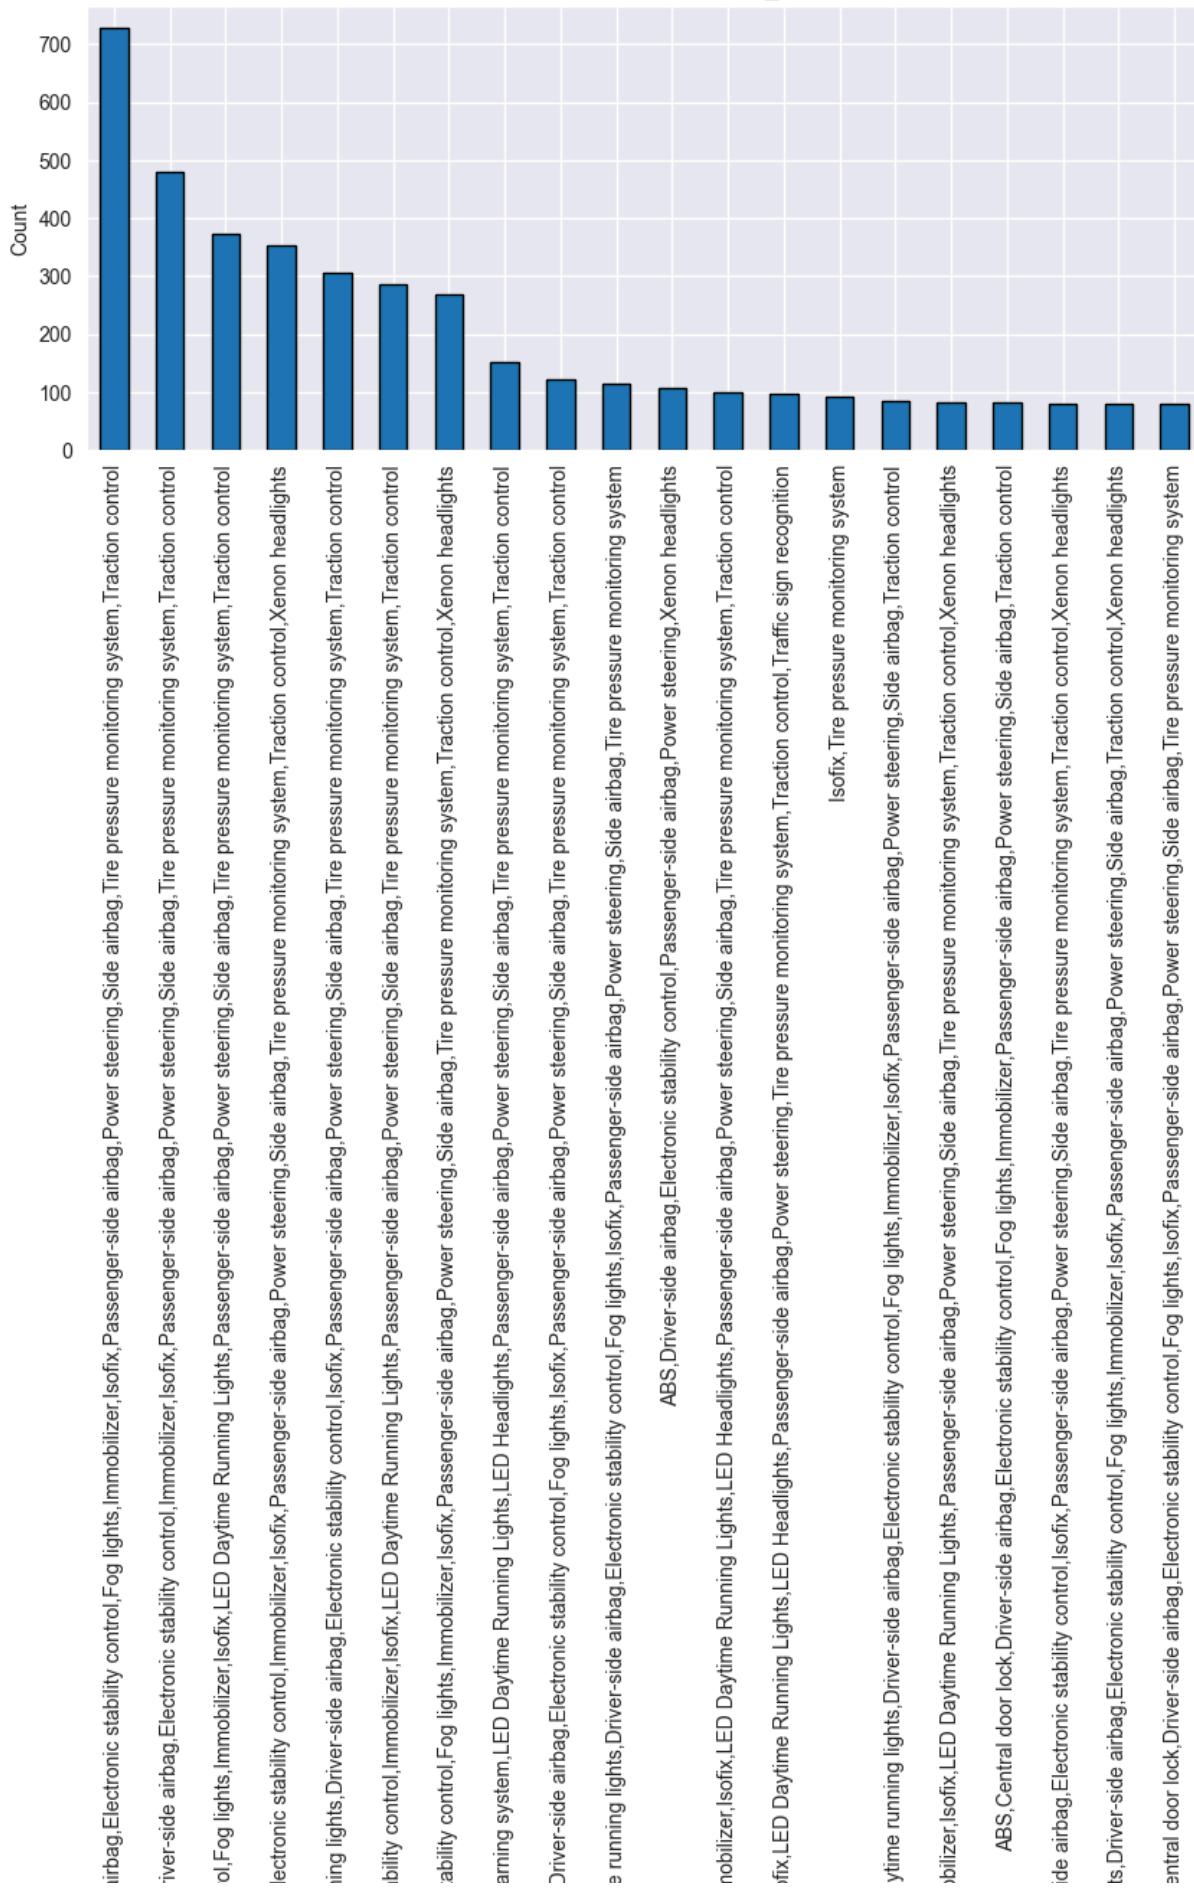


```
/var/folders/q9/_3w68hr55rdcl70s4bt4jjbh0000gn/T/ipykernel_66612/2463149345.py:28: UserWarning: Tight layout not applied. The bottom and top margins can not be made large enough to accommodate all Axes decorations.
plt.tight_layout()
```



```
/var/folders/q9/_3w68hr55rdcl70s4bt4jjbh0000gn/T/ipykernel_66612/2463149345.py:28: UserWarning: Tight layout not applied. The bottom and top margins can not be made large enough to accommodate all Axes decorations.  
plt.tight_layout()
```

Frequency Distribution of Safety_Security



ABS,C

ABS,Central door lock,Daytime running light|,Driver-side si

ABS,Central door lock,Daytime running lights,Di

ABS,Central door lock,Daytime running lights,Driver-side airbag,Electronic stability contr
ABS,Central door lock,Daytime running lights,Driver-side airbag,E|

ABS,Central door lock,Daytime runn

ABS,Central door lock,Daytime running lights,Driver-side airbag,Electronic sta

ABS,Central door lock,Daytime running lights,Driver-side airbag,Electronic st|
ABS,Adaptive headlights,Central door lock,Daytime running lights,Driver-side airbag,Electronic stability control,|
ABS,Blind spot monitor,Central door lock,Daytime running lights,Driver-side airbag,Electronic stability control,|
ABS,Central door lock,Daytime running lights,Driver-side airbag,Electronic stability control,Fog lights,Imm

ABS,Central door lock,Daytime running lights,Driver-side airbag,Electronic stability control,Fog lights,Imm

ABS,Central door lock,Daytime running lights,|,Immobilizer,|sofix,Lane departure w|

ABS,Central door lock,Daytime running lights,|

ABS,Central door lock,Daytime running lights,|,Immobilizer,|sofix,Lane departure w|

ABS,Central door lock,Daytime running lights,Driver-side airbag,Electronic stability control,|,Imm

ABS,Central door lock,Daytime running lights,Driver-side airbag,Electronic stability contr

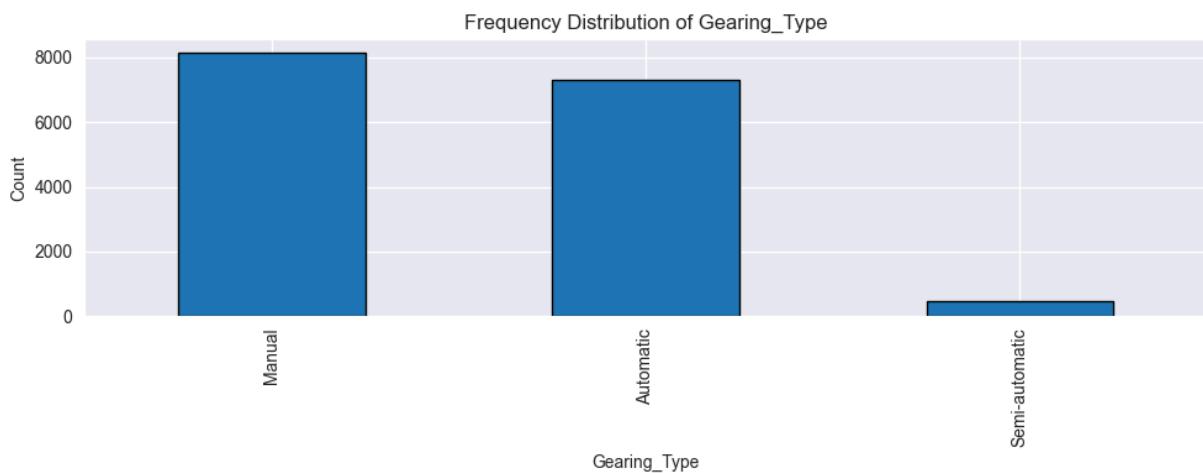
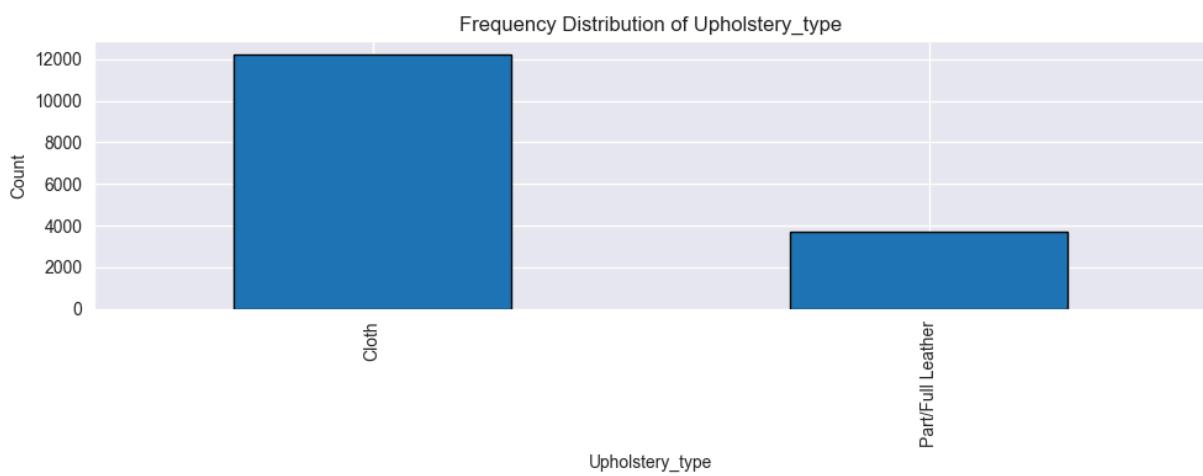
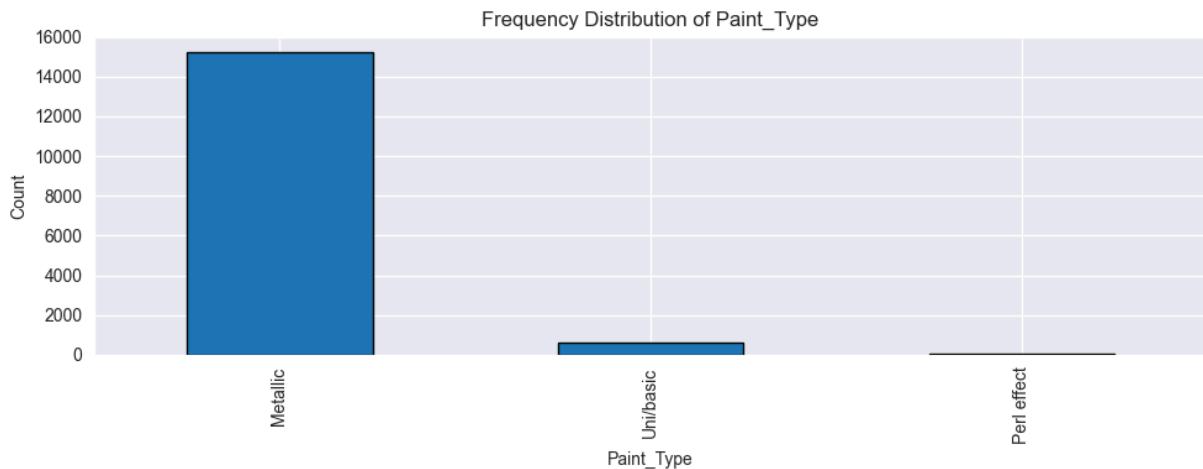
ABS,Central door lock,Daytime running lights,Driver-side airbag,E|

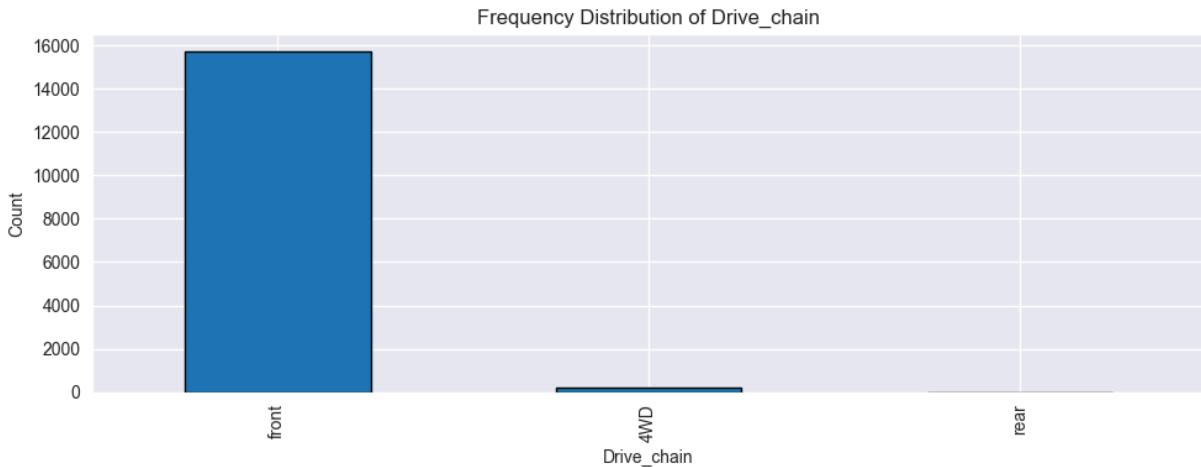
ABS,Central door lock,Daytime running lights,Driver-side airbag,Electronic sta

ABS,Central door lock,Daytime running lights,Driver-side airbag,E|

ABS,Central door lock,Daytime running lights,Driver-side airbag,Di

ABS,Central door lock,Daytime running lights,Driver-side airbag,Driver-side





Note: Look carefully at the values stored in columns `["Comfort_Convenience", "Entertainment_Media", "Extras", "Safety_Security"]`.

Should they be considered categorical? Should they be dropped or handled any other way?

2.1.4 [3 marks]

Fix columns with low frequency values and class imbalances.

Some information regarding values in the `Type` column that may help:

- 'Pre-registered' cars are ones which have already been registered previously by the seller.
- 'New' cars are not necessarily new cars, but new-like cars. These might also have multiple owners due to multiple pre-registrations as well.
- 'Employee's car' are cars used by employees over a short period of time and small distance.
- 'Demonstration' cars are used for trial purposes and also driven for a short time and distance.

Based on these, you can handle this particular column. For other columns, decide a strategy on your own.

In [126...]

```
# Fix columns as needed
# Fixing the 'Type' column based on domain logic
type_map = {
    "Pre-registered": "Nearly_New",
    "New": "Nearly_New",
    "Employee's car": "Nearly_New",
    "Demonstration": "Nearly_New"
}

df[["Type"]] = df[["Type"]].replace(type_map)

# Everything else becomes a generic used category
```

```

df["Type"] = df["Type"].fillna("Used_Regular")
df["Type"] = df["Type"].apply(lambda x: x if x in ["Nearly_New"] else "Used_Regular")

categorical_cols = df.select_dtypes(include=['object', 'category']).columns
categorical_cols = [col for col in categorical_cols]

# Function to merge rare categories (<1%) into 'Other'
def merge_rare_categories(series, threshold=0.01):
    freq = series.value_counts(normalize=True)
    rare = freq[freq < threshold].index
    return series.replace(rare, "Other")

# Apply globally
for col in categorical_cols:
    df[col] = merge_rare_categories(df[col])

df.head()

```

Out[126...]

	make_model	body_type	price	vat	km	Type	Fuel	Gears
0	Audi A1	Sedans	15770	VAT deductible	56013.0	Used-Regular	Diesel	7.0
1	Audi A1	Sedans	14500	Price negotiable	80000.0	Used-Regular	Benzine	7.0
2	Audi A1	Sedans	14640	VAT deductible	83450.0	Used-Regular	Diesel	7.0
3	Audi A1	Sedans	14500	VAT deductible	73000.0	Used-Regular	Diesel	6.0
4	Audi A1	Sedans	16790	VAT deductible	16200.0	Used-Regular	Diesel	7.0

5 rows × 23 columns

2.1.5 [3 marks]

Identify target variable and plot the frequency distributions. Apply necessary transformations.

In [127...]

```

# Plot histograms for target feature
import matplotlib.pyplot as plt

target = "price" # Explicitly chosen based on dataset + project objective
print("Selected target variable:", target)

```

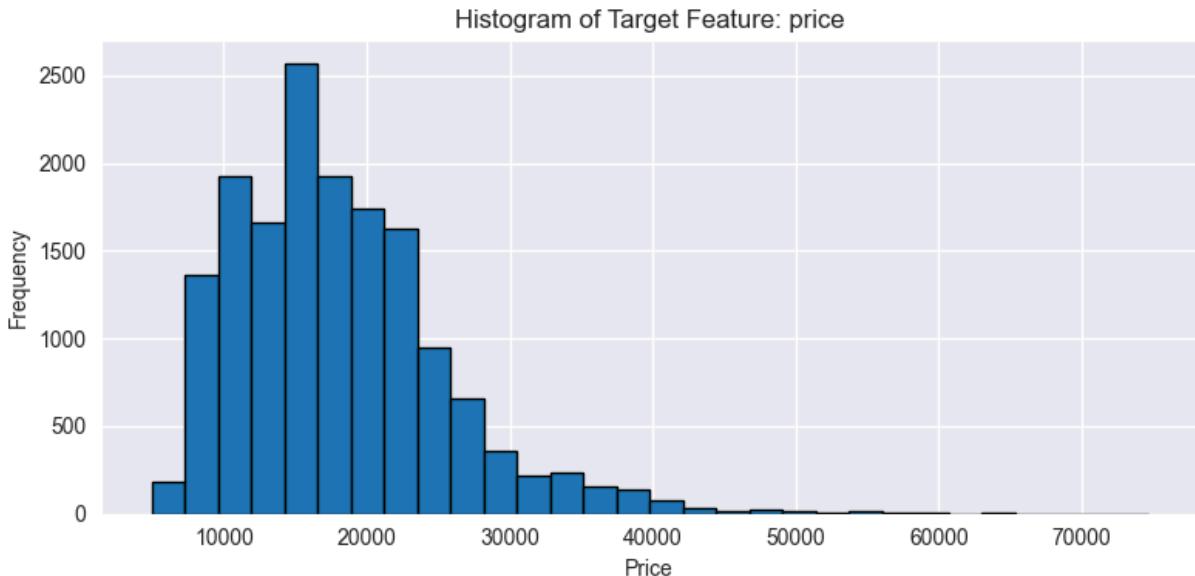
```

if df[target].dtype == "object":
    plt.figure(figsize=(8,4))
    df[target].value_counts().plot(kind="bar", edgecolor="black")
    plt.title(f"Frequency Distribution of Target: {target}")
    plt.xlabel(target)
    plt.ylabel("Count")
    plt.tight_layout()
    plt.show()

else:
    plt.figure(figsize=(8,4))
    plt.hist(df[target], bins=30, edgecolor="black")
    plt.title(f"Histogram of Target Feature: {target}")
    plt.xlabel("Price")
    plt.ylabel("Frequency")
    plt.tight_layout()
    plt.show()

```

Selected target variable: price



The target variable seems to be skewed. Perform suitable transformation on the target.

In [128...]: # Transform the target feature

```

import matplotlib.pyplot as plt
import seaborn as sns

plt.figure(figsize=(8,4))
sns.histplot(df["price"], bins=30, kde=True)
plt.title("Price Distribution (Before Transformation)")
plt.show()

```

```

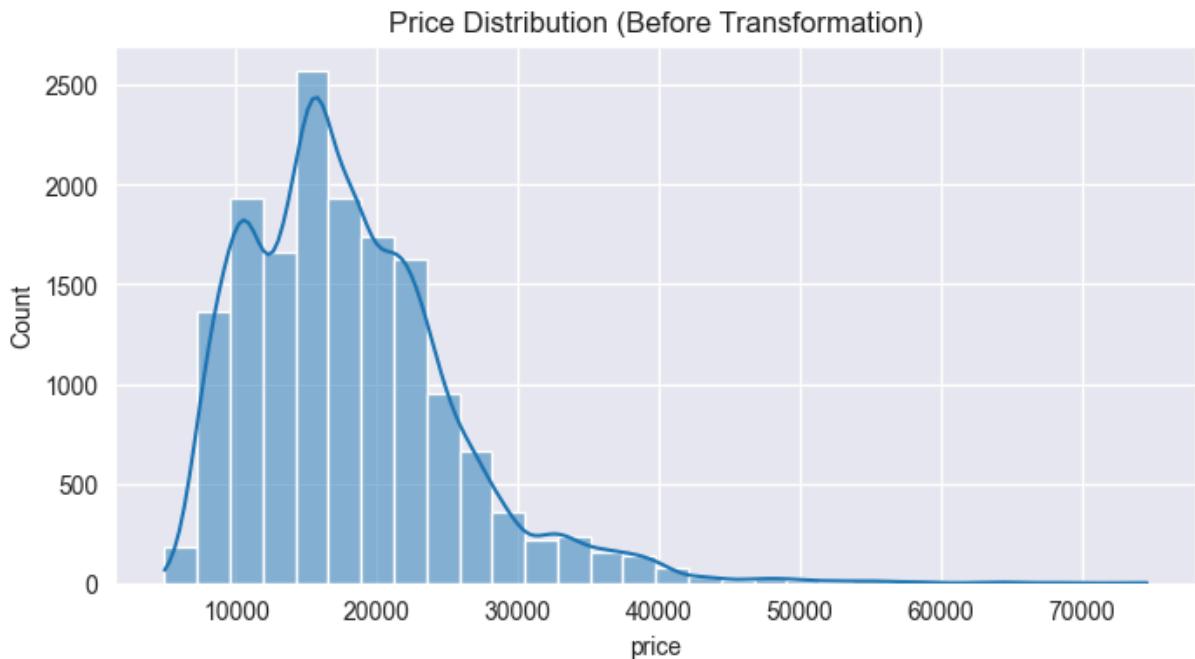
print("Skewness before transformation:", df["price"].skew())

# Apply log transformation
df["price_log"] = np.log1p(df["price"])

print("Skewness after log transformation:", df["price_log"].skew())

plt.figure(figsize=(8,4))
sns.histplot(df["price_log"], bins=30, kde=True)
plt.title("Price Distribution After Log Transformation")
plt.show()

```



Skewness before transformation: 1.236169412899669

Skewness after log transformation: -0.0314736417467197



2.2 Correlation analysis [6 marks]

2.2.1 [3 marks]

Plot the correlation map between features and target variable.

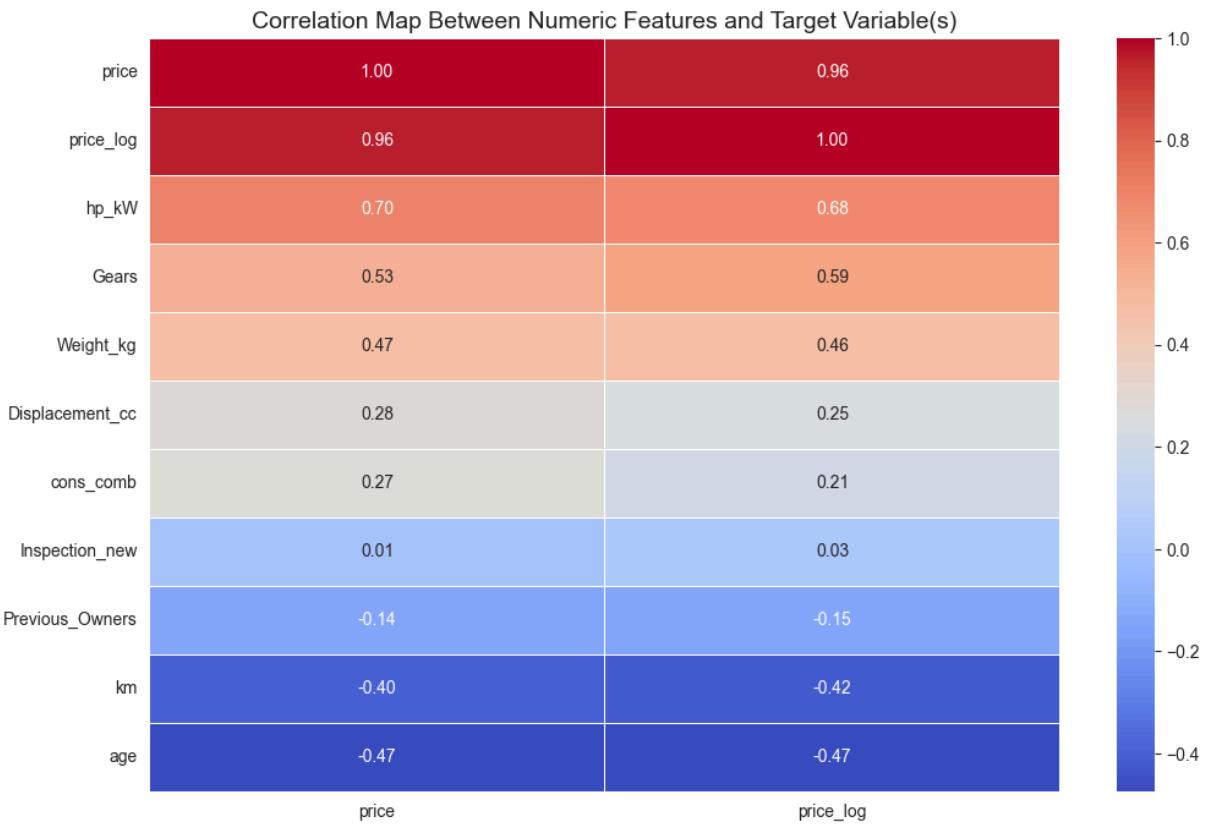
```
In [129...]: # Visualise correlation
import seaborn as sns
import matplotlib.pyplot as plt
import seaborn as sns
import matplotlib.pyplot as plt

# Select numeric columns
numeric_df = df.select_dtypes(include=['int64', 'float64'])

# Include both original and transformed target
targets = ["price"]
if "price_log" in df.columns:
    targets.append("price_log")

corr_matrix = numeric_df[targets + [col for col in numeric_df.columns if col

plt.figure(figsize=(12, 8))
sns.heatmap(
    corr_matrix[[target for target in targets]].sort_values(by=targets, ascending=False),
    annot=True,
    cmap="coolwarm",
    fmt=".2f",
    linewidths=.5
)
plt.title("Correlation Map Between Numeric Features and Target Variable(s)", fontweight="bold")
plt.show()
```



2.2.2 [3 marks]

Analyse correlation between categorical features and target variable.

```
In [130]: # Comparing average values of target for different categories
category_analysis = {}

for col in categorical_cols:
    summary = df.groupby(col)[["price", "price_log"]].agg(
        avg_price = ("price", "mean"),
        avg_price_log = ("price_log", "mean"),
        count = ("price", "count")
    ).sort_values("avg_price", ascending=False)

    category_analysis[col] = summary
    print(f"\n===== {col} =====")
    print(summary.head(10))
    # show top 10 categories

import matplotlib.pyplot as plt

for col in categorical_cols:
    if df[col].nunique() <= 15: # avoid exploding plots
        plt.figure(figsize=(10,4))
        df.groupby(col)["price"].mean().sort_values().plot(kind="bar", edgecolor="black")
        plt.title(f"Average Price for Each {col} Category")
        plt.xlabel(col)
        plt.ylabel("Average Price")
```

```
plt.tight_layout()  
plt.show()
```

===== make_model =====

make_model	avg_price	avg_price_log	count
Renault Espace	30080.211907	10.271724	991
Opel Insignia	21463.451886	9.912658	2598
Audi A3	20996.693252	9.928165	3097
Audi A1	18864.688982	9.817643	2614
Opel Astra	15840.834059	9.626041	2525
Other	13657.885714	9.504893	35
Renault Clio	11940.320827	9.335822	1839
Opel Corsa	11061.841606	9.276059	2216

===== body_type =====

body_type	avg_price	avg_price_log	count
Van	30789.209302	10.294540	817
Station wagon	18542.296981	9.753456	3677
Sedans	17699.040480	9.716144	8004
Other	17023.604520	9.622932	177
Compact	15076.206481	9.555735	3240

===== vat =====

vat	avg_price	avg_price_log	count
VAT deductible	18185.887131	9.730286	15044
Price negotiable	15234.823192	9.556299	871

===== Type =====

Type	avg_price	avg_price_log	count
Nearly_New	22127.53278	9.928387	4820
Used-Regular	16241.84849	9.630566	11095

===== Fuel =====

Fuel	avg_price	avg_price_log	count
Diesel	18177.407783	9.734926	7298
Benzine	17899.681329	9.709328	8548
Other	17287.231884	9.639463	69

===== Comfort_Convenience =====

Comfort_Convenience	avg_price \
Air conditioning,Armrest,Automatic climate control...	23292.099379
Air conditioning,Automatic climate control,Cruise...	20244.243750
Other	18172.657168
Air conditioning,Armrest,Automatic climate control...	17637.422460
Air conditioning,Electrical side mirrors,Hill Hold...	9370.036082

Comfort_Convenience	avg_price_log	count
Air conditioning,Armrest,Automatic climate control...	10.017481	161
Air conditioning,Automatic climate control,Cruise...	9.898748	160
Other	9.730398	15019
Air conditioning,Armrest,Automatic climate control...	9.762393	187
Air conditioning,Electrical side mirrors,Hill Hold...	9.131270	388

===== Entertainment_Media =====

	avg_price	\
Entertainment_Media		
Bluetooth,Digital radio,Hands-free equipment,On...	23972.817544	
Bluetooth,Hands-free equipment,On-board compute...	22785.580583	
Bluetooth,Digital radio,Hands-free equipment,MP...	21847.373016	
Bluetooth,Radio	20092.339483	
Bluetooth,CD player,Digital radio,Hands-free eq...	19290.977143	
Other	19123.484680	
Bluetooth,CD player,Hands-free equipment,MP3,On...	19100.259939	
Bluetooth,CD player,Hands-free equipment,MP3,On...	17878.138614	
Bluetooth,Hands-free equipment,On-board compute...	17709.803191	
Radio	17629.154122	

avg_price_log count

Entertainment_Media		
Bluetooth,Digital radio,Hands-free equipment,On...	10.018384	285
Bluetooth,Hands-free equipment,On-board compute...	9.963690	515
Bluetooth,Digital radio,Hands-free equipment,MP...	9.911307	252
Bluetooth,Radio	9.857507	271
Bluetooth,CD player,Digital radio,Hands-free eq...	9.784390	175
Other	9.770708	5940
Bluetooth,CD player,Hands-free equipment,MP3,On...	9.819240	327
Bluetooth,CD player,Hands-free equipment,MP3,On...	9.742608	404
Bluetooth,Hands-free equipment,On-board compute...	9.743829	188
Radio	9.701190	558

===== Extras =====

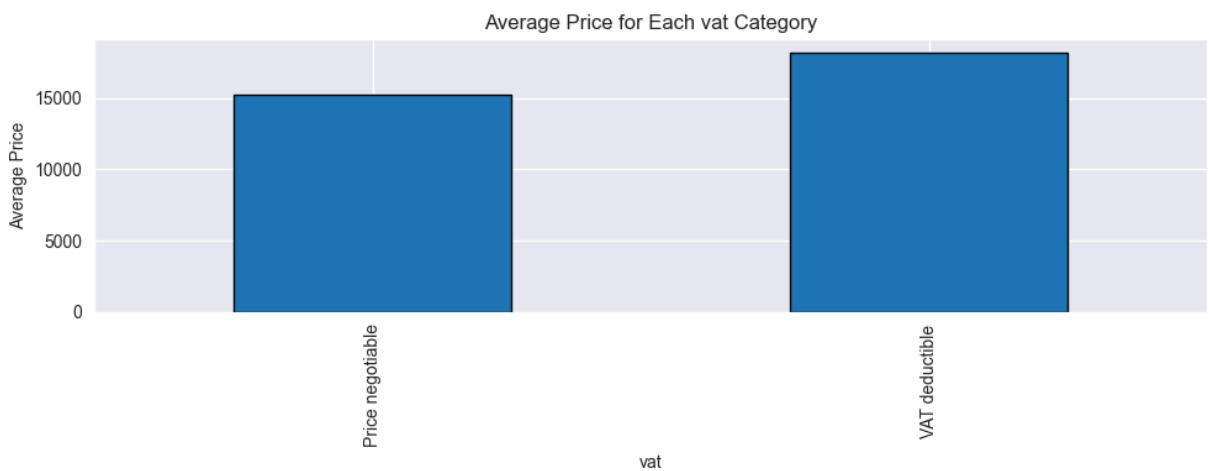
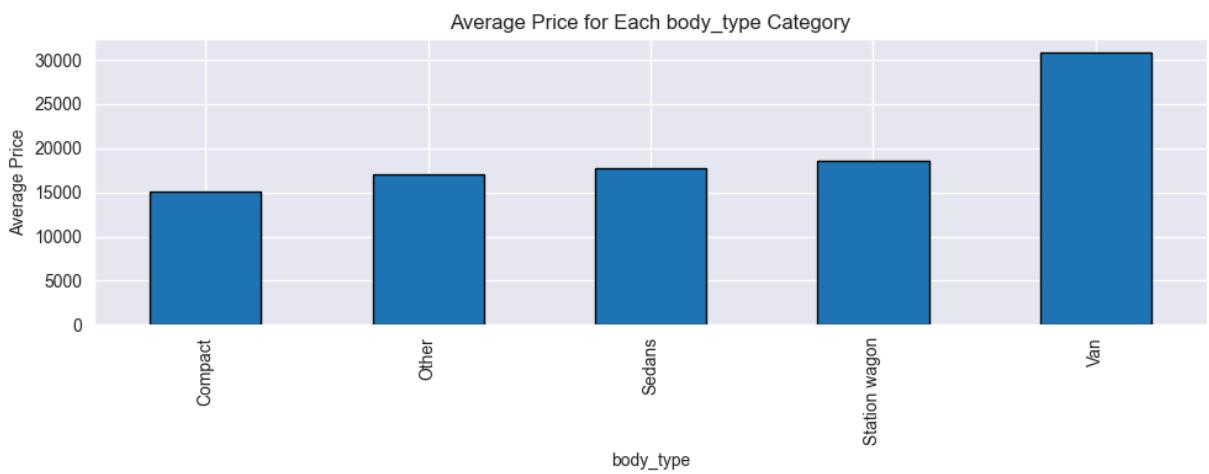
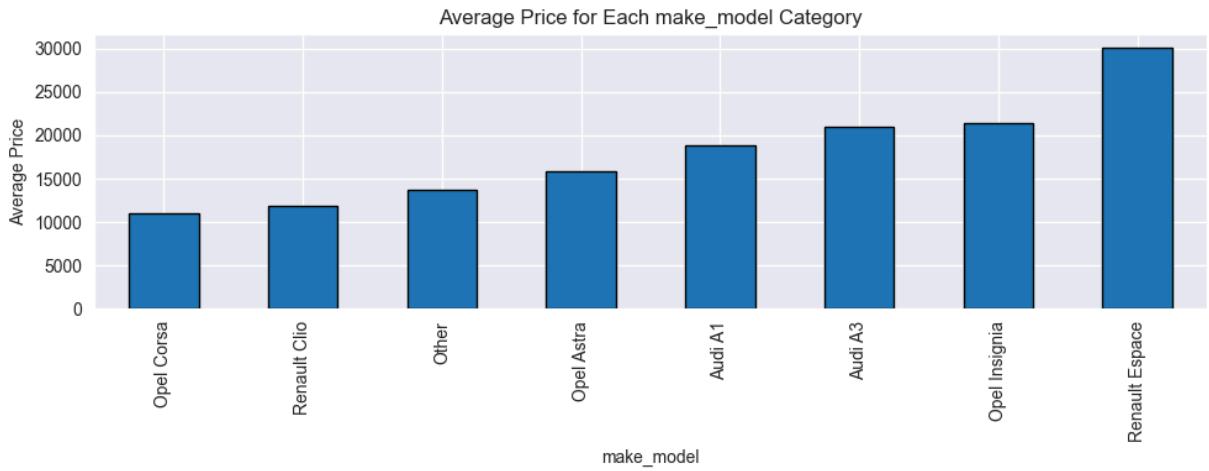
	avg_price	\
Extras		
Alloy wheels,Sport suspension	21285.880682	
Alloy wheels,Roof rack,Touch screen,Voice Control	20755.177083	
Alloy wheels,Catalytic Converter,Voice Control	20631.334975	
Alloy wheels,Touch screen	20254.259684	
Alloy wheels,Sport seats,Sport suspension	20104.918803	
Alloy wheels,Touch screen,Voice Control	19978.365809	
Alloy wheels,Roof rack	19893.614367	
Alloy wheels,Catalytic Converter,Touch screen	19619.412281	
Alloy wheels,Voice Control	19382.783505	
Other	19341.330786	

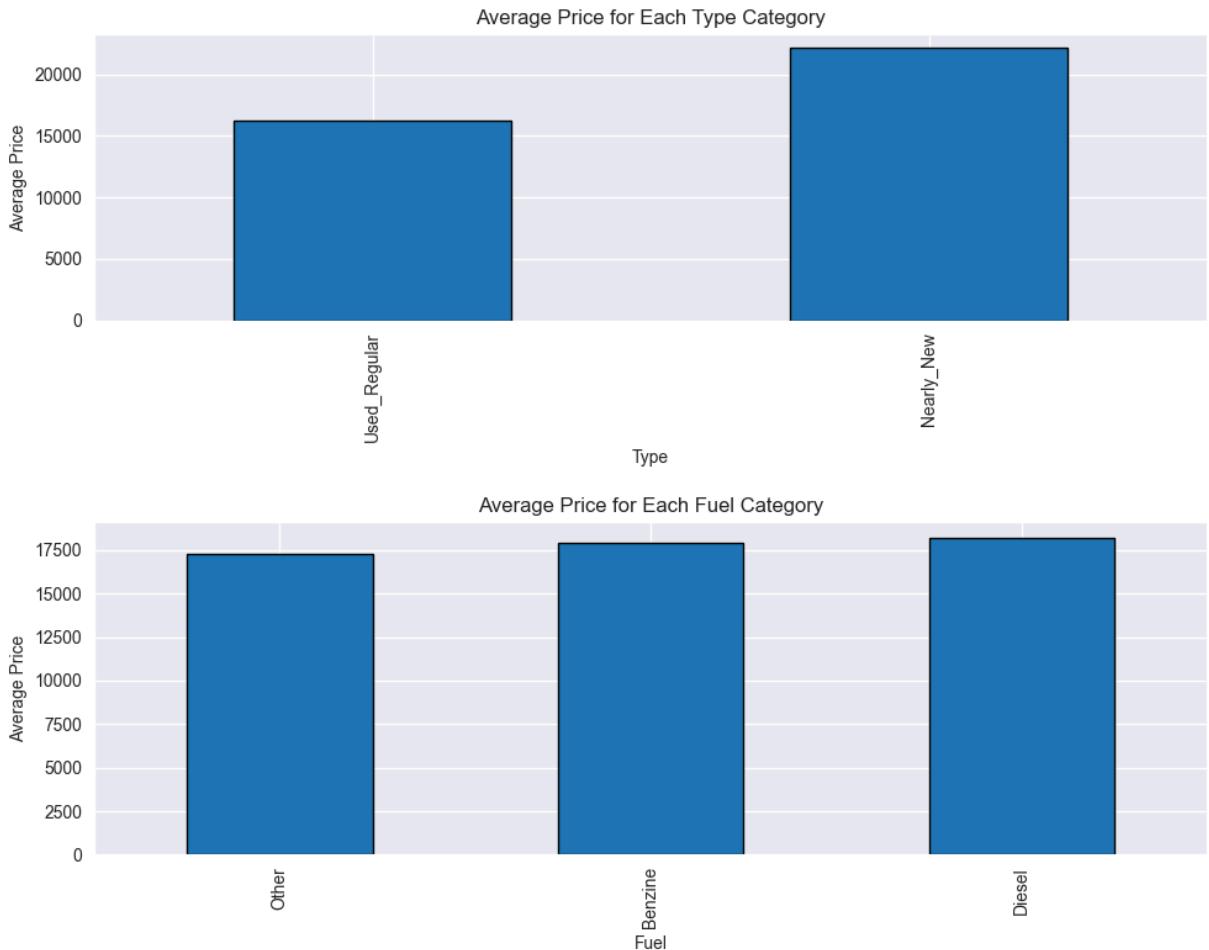
avg_price_log count

Extras		
Alloy wheels,Sport suspension	9.942534	176
Alloy wheels,Roof rack,Touch screen,Voice Control	9.896251	192
Alloy wheels,Catalytic Converter,Voice Control	9.892297	203
Alloy wheels,Touch screen	9.825623	697
Alloy wheels,Sport seats,Sport suspension	9.870663	234
Alloy wheels,Touch screen,Voice Control	9.831347	544
Alloy wheels,Roof rack	9.836694	529
Alloy wheels,Catalytic Converter,Touch screen	9.813442	228
Alloy wheels,Voice Control	9.785768	582
Other	9.807676	4710

===== Safety_Security =====

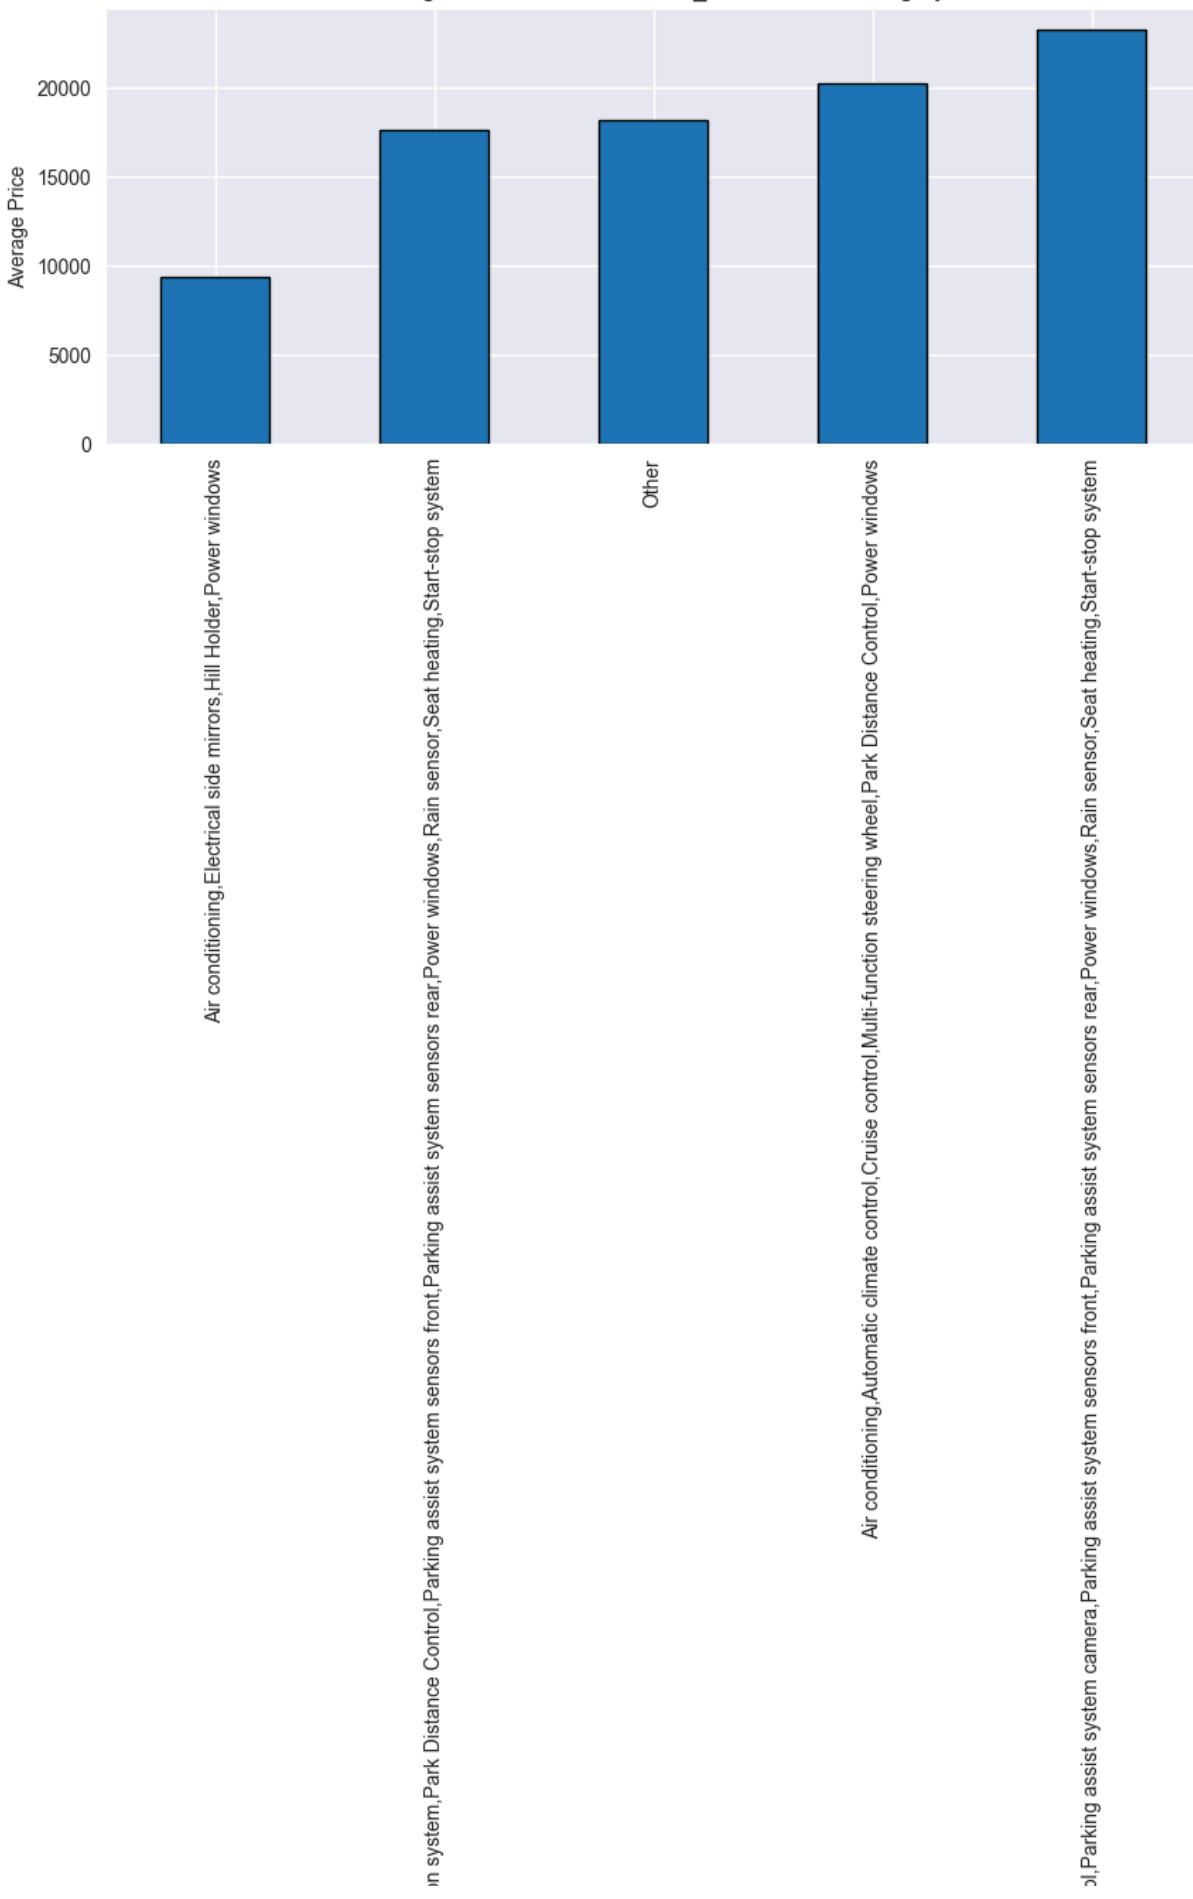
		avg_price	\
Safety_Security			
ABS,Central door lock,Daytime running lights,Dr...	20902.099432		
Other	18689.712729		
ABS,Central door lock,Daytime running lights,Dr...	17821.223048		
ABS,Central door lock,Daytime running lights,Dr...	15101.008043		
ABS,Central door lock,Daytime running lights,Dr...	14085.145405		
ABS,Central door lock,Daytime running lights,Dr...	13811.362745		
ABS,Central door lock,Daytime running lights,Dr...	12508.612500		
ABS,Central door lock,Daytime running lights,Dr...	11770.559441		
		avg_price_log	count
Safety_Security			
ABS,Central door lock,Daytime running lights,Dr...	9.930039	352	
Other	9.755765	13120	
ABS,Central door lock,Daytime running lights,Dr...	9.774915	269	
ABS,Central door lock,Daytime running lights,Dr...	9.596806	373	
ABS,Central door lock,Daytime running lights,Dr...	9.510400	729	
ABS,Central door lock,Daytime running lights,Dr...	9.448312	306	
ABS,Central door lock,Daytime running lights,Dr...	9.394061	480	
ABS,Central door lock,Daytime running lights,Dr...	9.344273	286	
===== Paint_Type =====			
	avg_price	avg_price_log	count
Paint_Type			
Metallic	18095.052866	9.725404	15246
Uni/basic	16745.040816	9.638642	637
Other	9820.312500	9.144565	32
===== Upholstery_type =====			
	avg_price	avg_price_log	count
Upholstery_type			
Part/Full Leather	23226.196142	9.981383	3681
Cloth	16459.243829	9.642348	12234
===== Gearing_Type =====			
	avg_price	avg_price_log	count
Gearing_Type			
Semi-automatic	23236.562900	9.972004	469
Automatic	21163.176237	9.906177	7297
Manual	14913.777396	9.540276	8149
===== Drive_chain =====			
	avg_price	avg_price_log	count
Drive_chain			
4WD	28261.416667	10.133491	204
front	17892.109123	9.715432	15707
Other	15332.500000	9.606166	4





```
/var/folders/q9/_3w68hr55rdcl70s4bt4jjbh0000gn/T/ipykernel_66612/1065706893.py:26: UserWarning: Tight layout not applied. The bottom and top margins can not be made large enough to accommodate all Axes decorations.  
plt.tight_layout()
```

Average Price for Each Comfort_Convenience Category



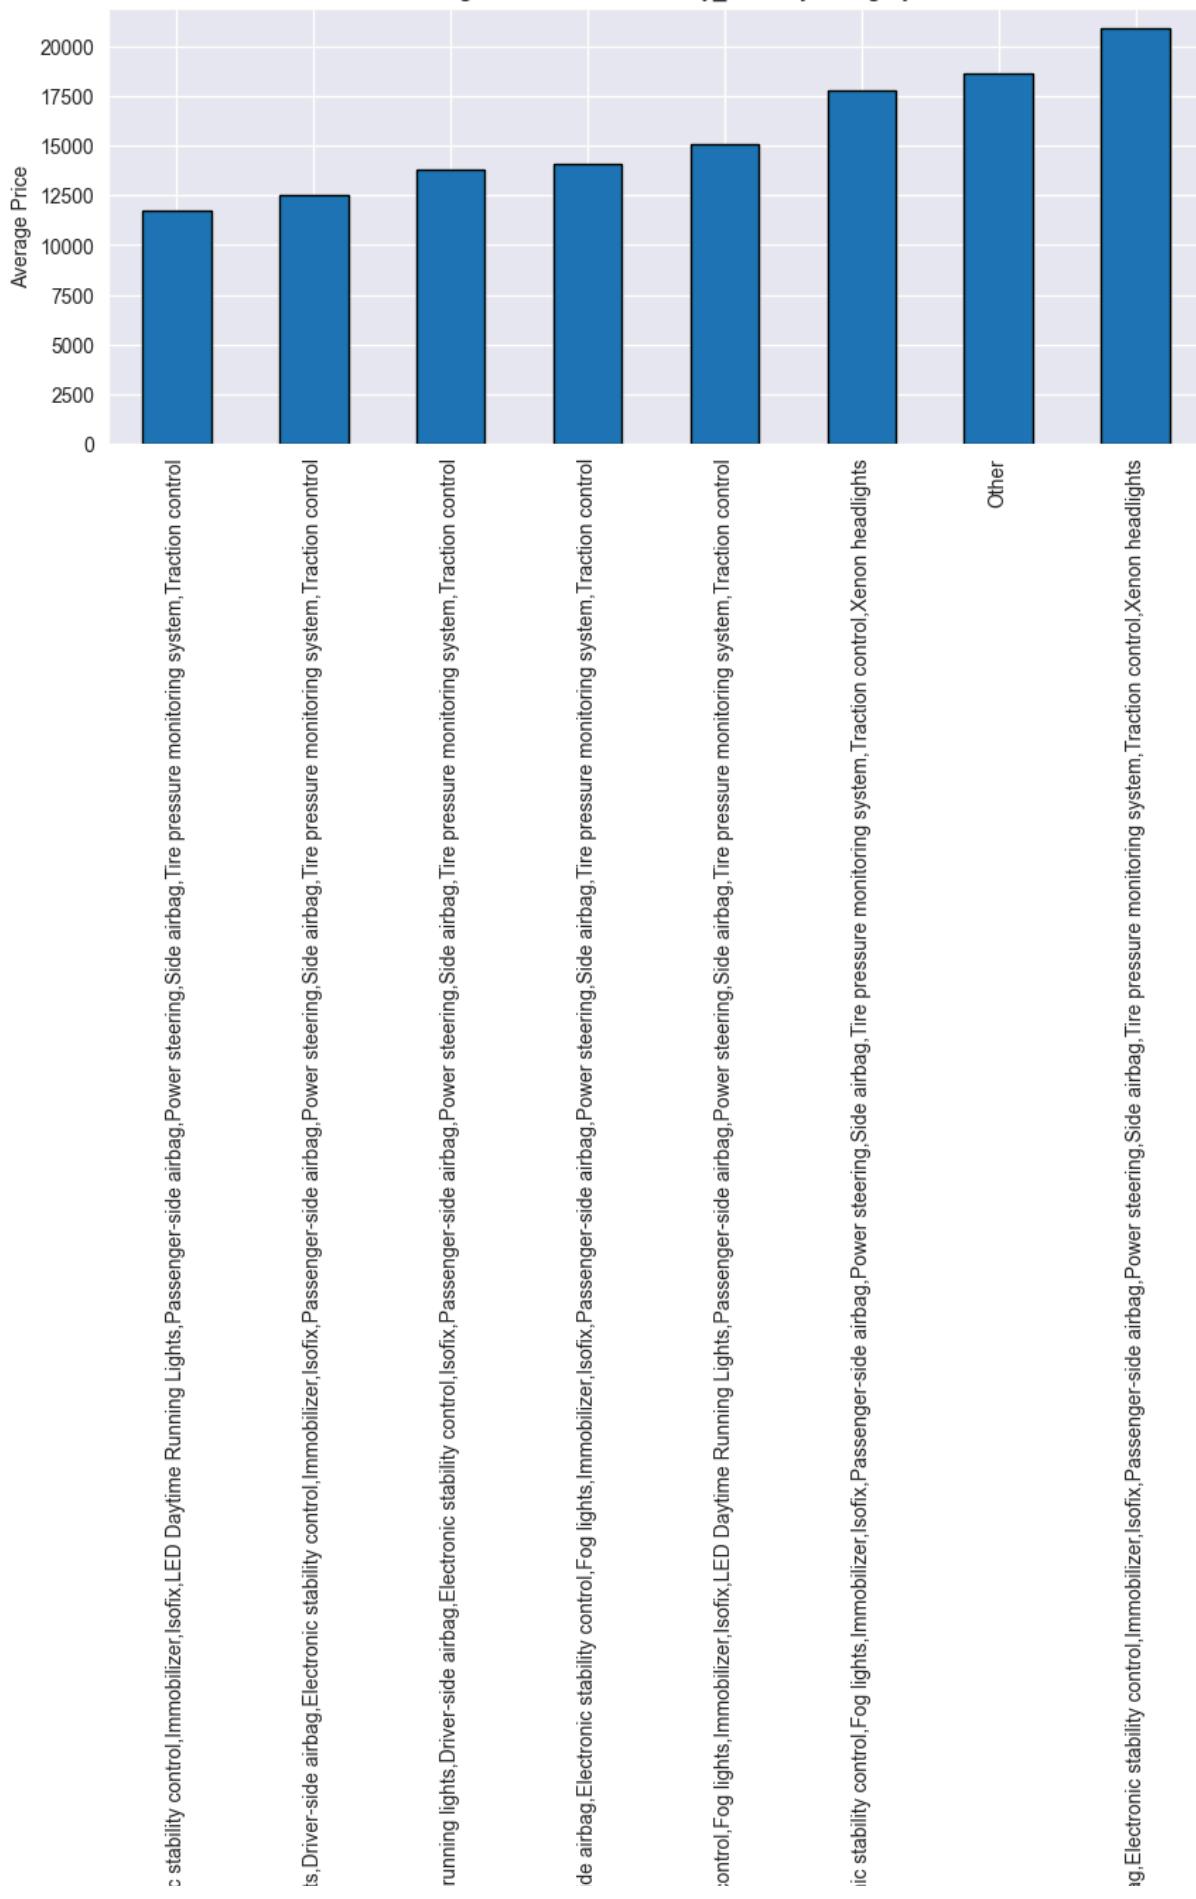
;Air conditioning,Armrest,Automatic climate control,Cruise control,Electrical side mirrors,Leather steering wheel,Light sensor,Lumbar support,Multi-function steering wheel,Navigation system,Park Distance Contr

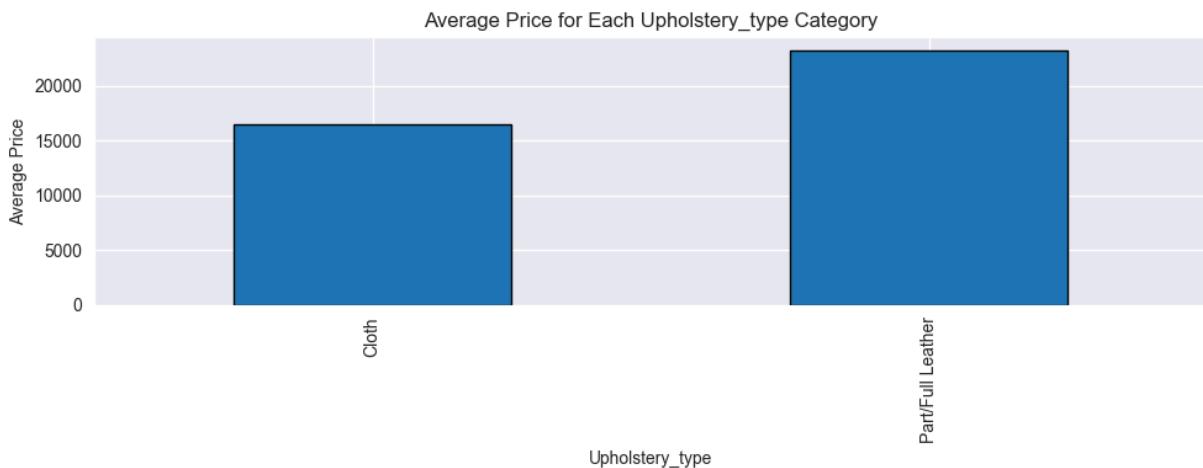
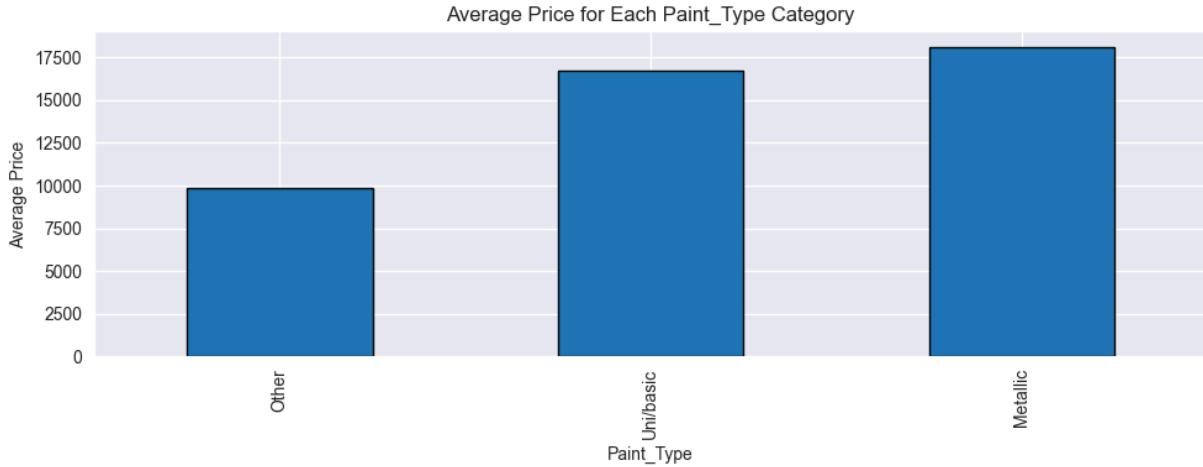
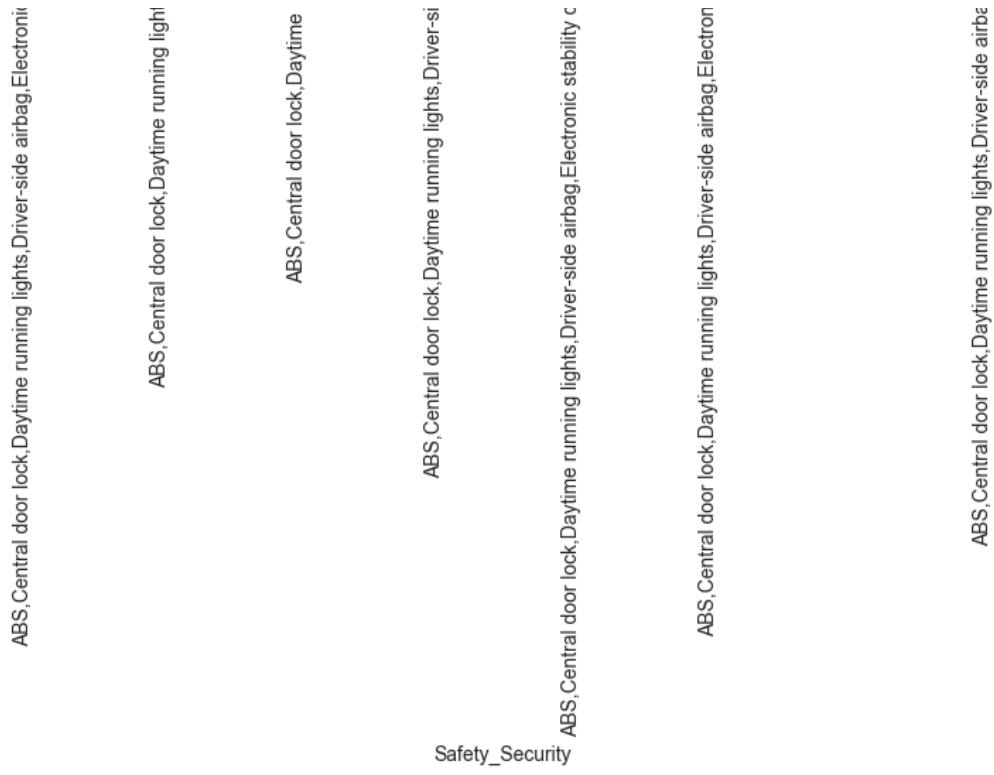
Air conditioning,Armrest,Automatic climate control,Cruise control,Electrical side mirrors,Leather steering wheel,Light sensor,Lumbar support,Multi-function steering wheel,Navigation system,Park Distance Contr

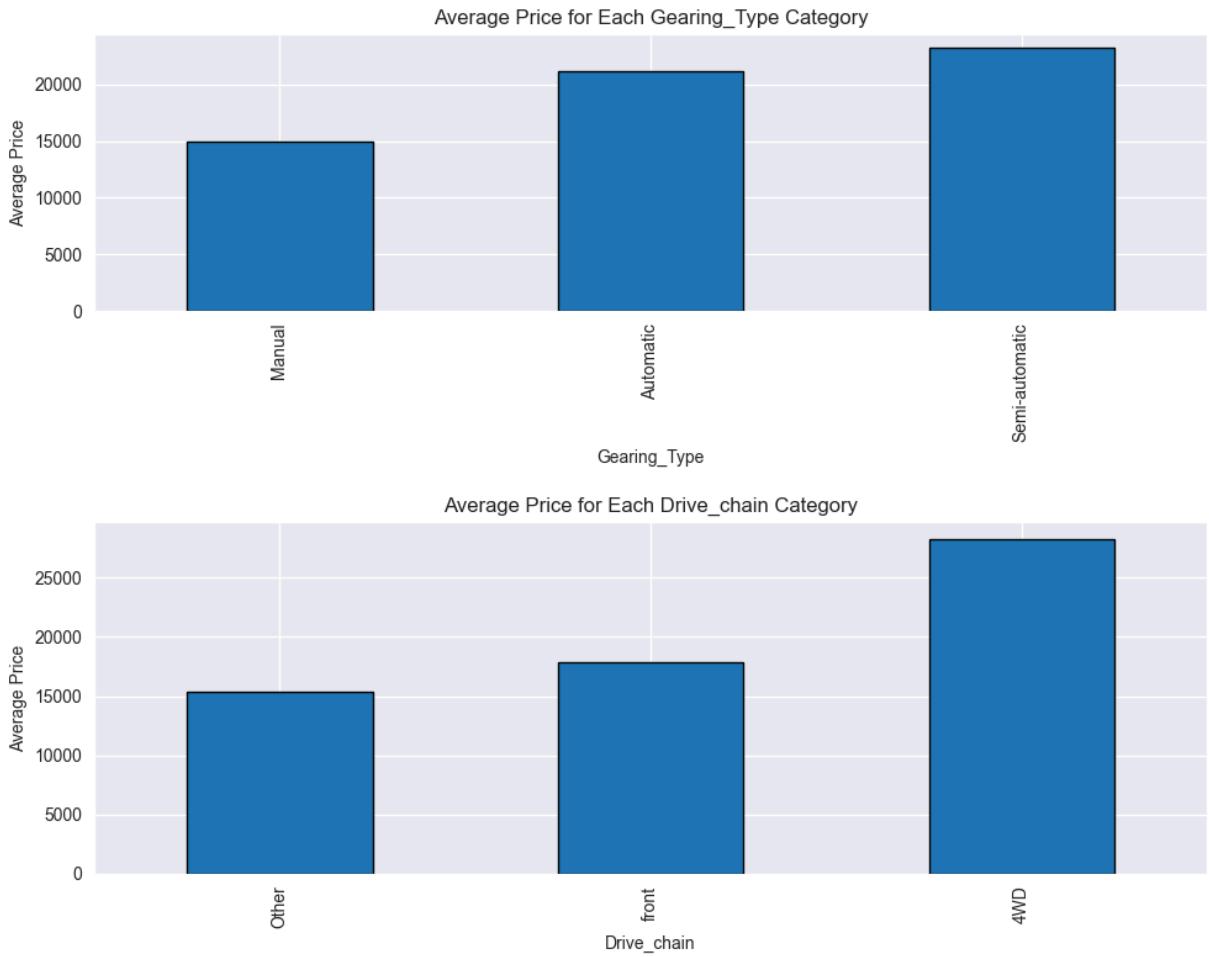
Comfort_Convenience

```
/var/folders/q9/_3w68hr55rdcl70s4bt4jjbh0000gn/T/ipykernel_66612/1065706893.py:26: UserWarning: Tight layout not applied. The bottom and top margins can not be made large enough to accommodate all Axes decorations.  
 plt.tight_layout()
```

Average Price for Each Safety_Security Category







2.3 Outlier analysis [5 marks]

2.3.1 [2 marks]

Identify potential outliers in the data.

```
In [131]: # Outliers present in each column
numeric_cols = df.select_dtypes(include=['int64', 'float64']).columns
numeric_cols
outlier_summary = {}

for col in numeric_cols:
    Q1 = df[col].quantile(0.25)
    Q3 = df[col].quantile(0.75)
    IQR = Q3 - Q1
    lower = Q1 - 1.5 * IQR
    upper = Q3 + 1.5 * IQR

    outliers = df[(df[col] < lower) | (df[col] > upper)][col]
    outlier_summary[col] = len(outliers)

# Print summary
print("Outlier counts per numeric column:")
for col, count in outlier_summary.items():
    print(f"{col}: {count}")
```

```

import matplotlib.pyplot as plt
import seaborn as sns

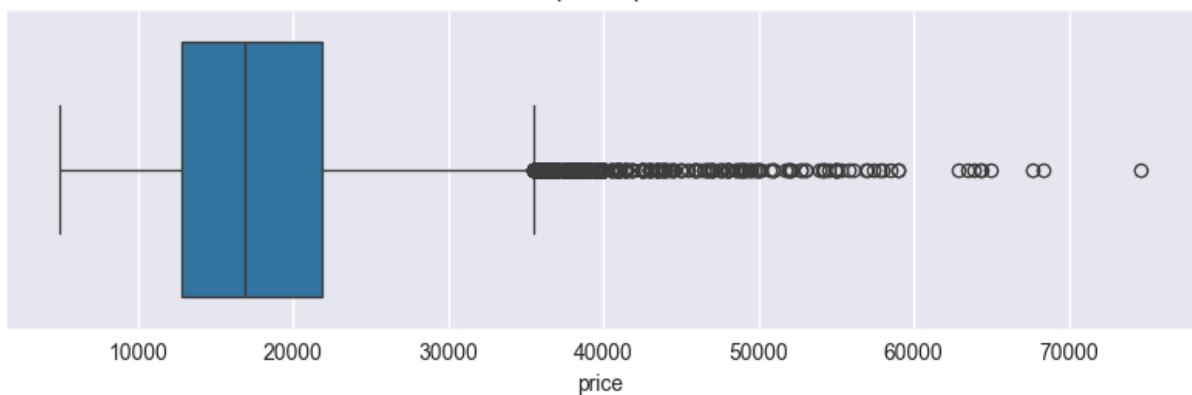
for col in numeric_cols:
    plt.figure(figsize=(8, 3))
    sns.boxplot(x=df[col])
    plt.title(f"Boxplot of {col}")
    plt.tight_layout()
    plt.show()

```

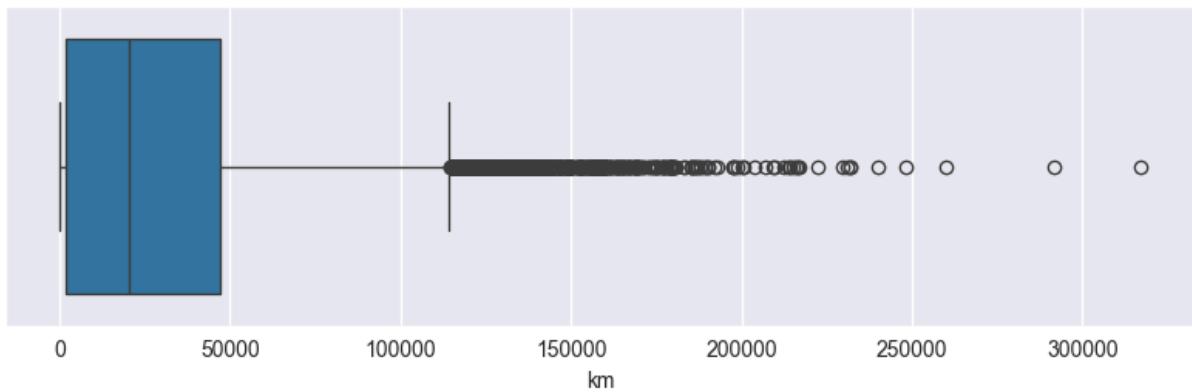
Outlier counts per numeric column:

price: 479
 km: 689
 Gears: 225
 age: 0
 Previous_Owners: 1757
 hp_kw: 361
 Inspection_new: 3932
 Displacement_cc: 21
 Weight_kg: 87
 cons_comb: 125
 price_log: 71

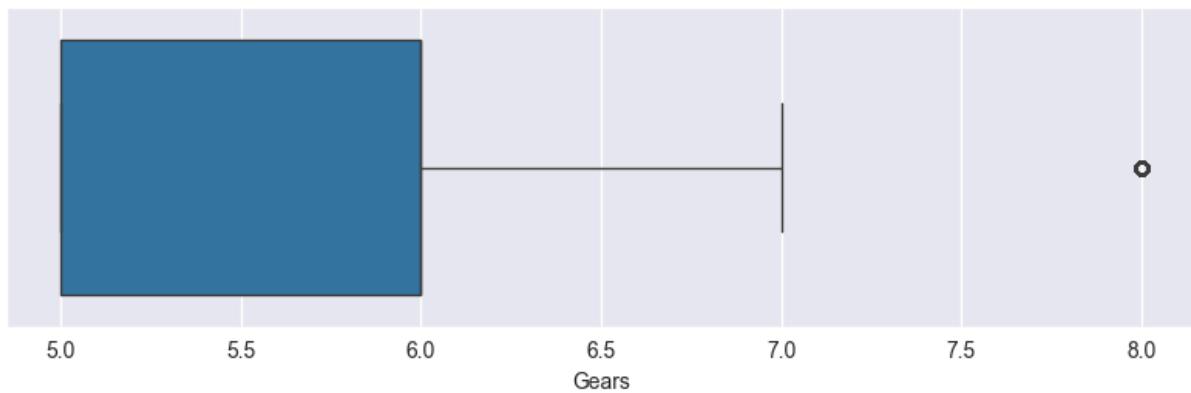
Boxplot of price



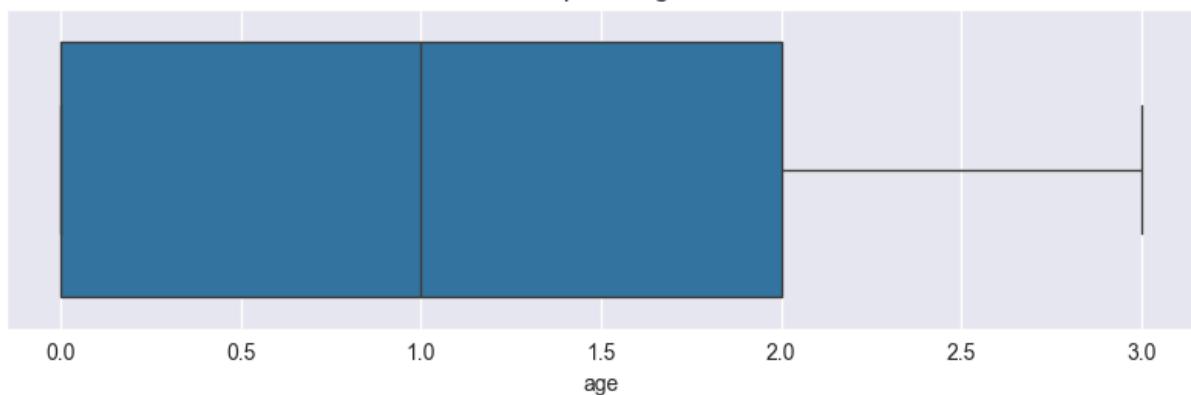
Boxplot of km



Boxplot of Gears



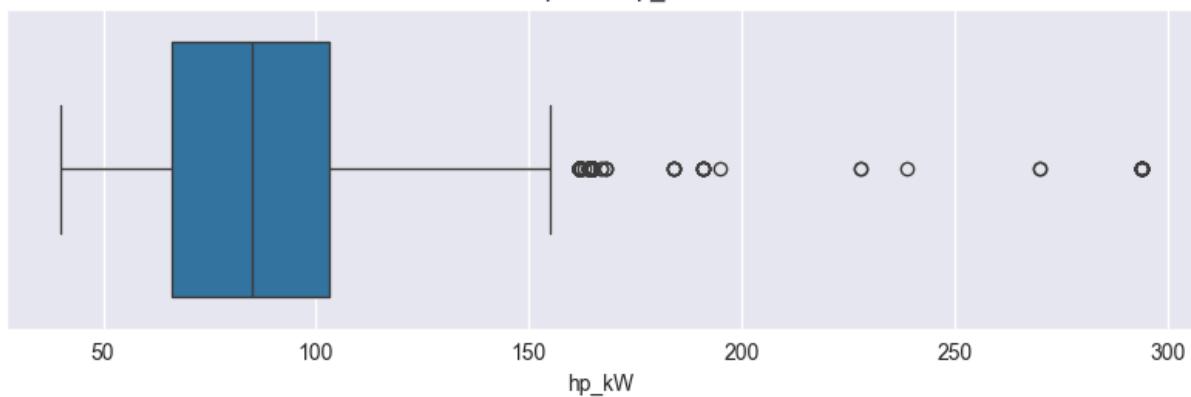
Boxplot of age



Boxplot of Previous_Owners



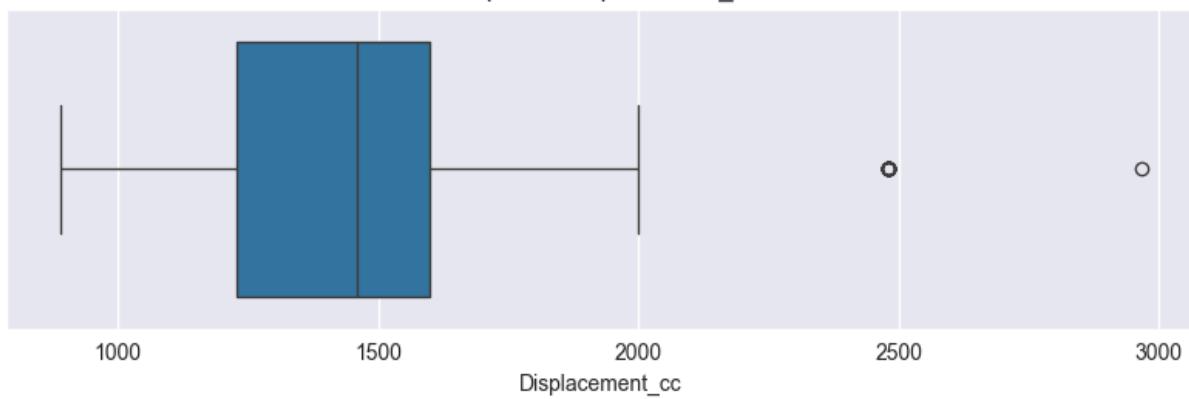
Boxplot of hp_kW



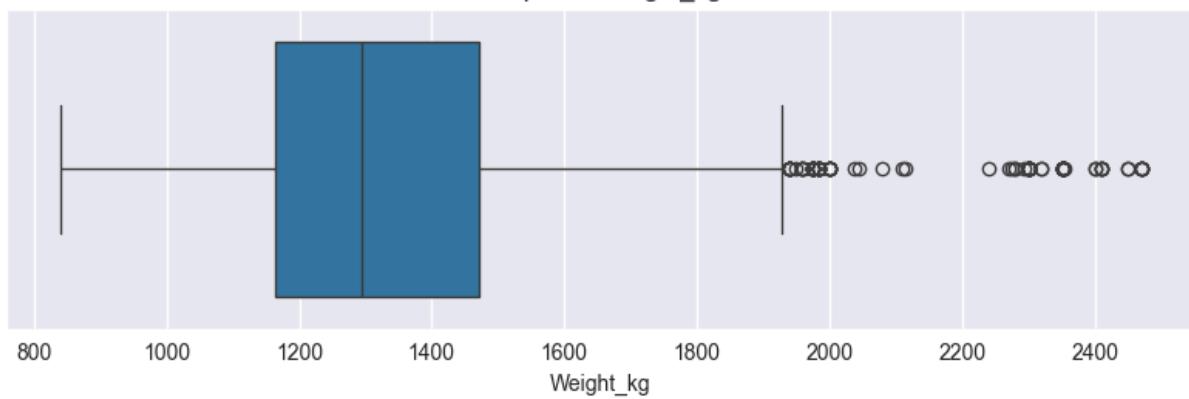
Boxplot of Inspection_new



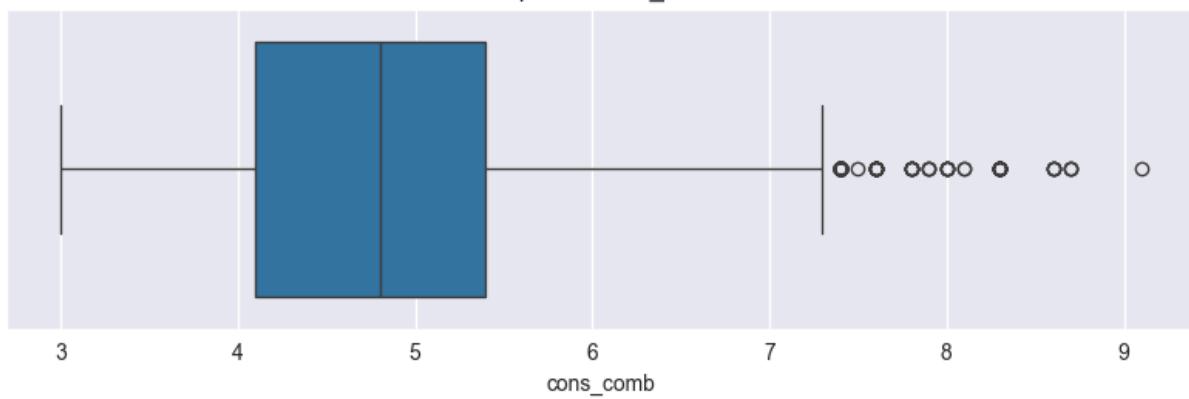
Boxplot of Displacement_cc

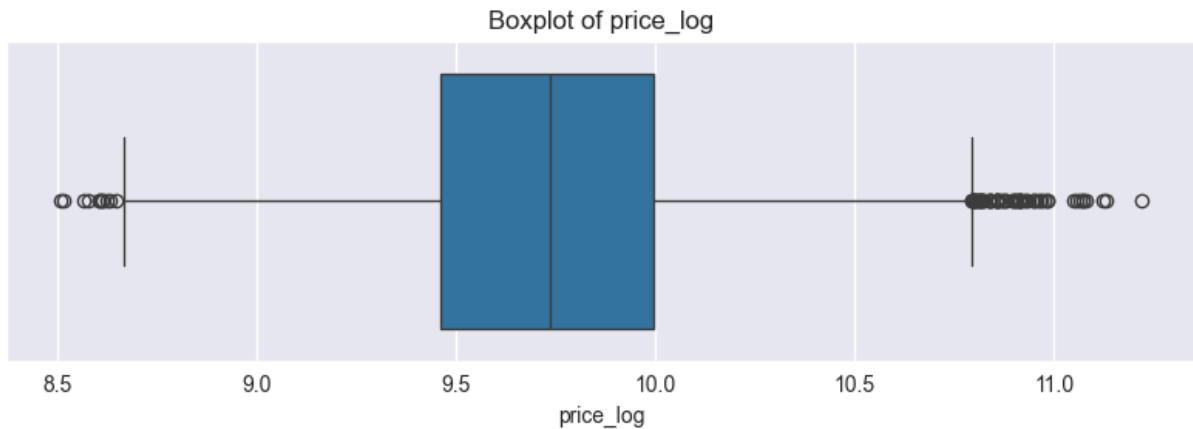


Boxplot of Weight_kg



Boxplot of cons_comb





2.3.2 [3 marks]

Handle the outliers suitably.

```
In [132]: # Handle outliers
numeric_cols = df.select_dtypes(include=['int64', 'float64']).columns.tolist()
numeric_cols

def winsorize(series, lower=0.01, upper=0.99):
    lower_bound = series.quantile(lower)
    upper_bound = series.quantile(upper)
    return series.clip(lower_bound, upper_bound)

numeric_cols_to_cap = [col for col in numeric_cols if col not in ["price_log"]]

for col in numeric_cols_to_cap:
    df[col] = winsorize(df[col])

from scipy.stats import zscore

outliers_post = (abs(df[numeric_cols].apply(zscore)) > 3).sum()
print("\nOutliers remaining after winsorization:")
print(outliers_post)

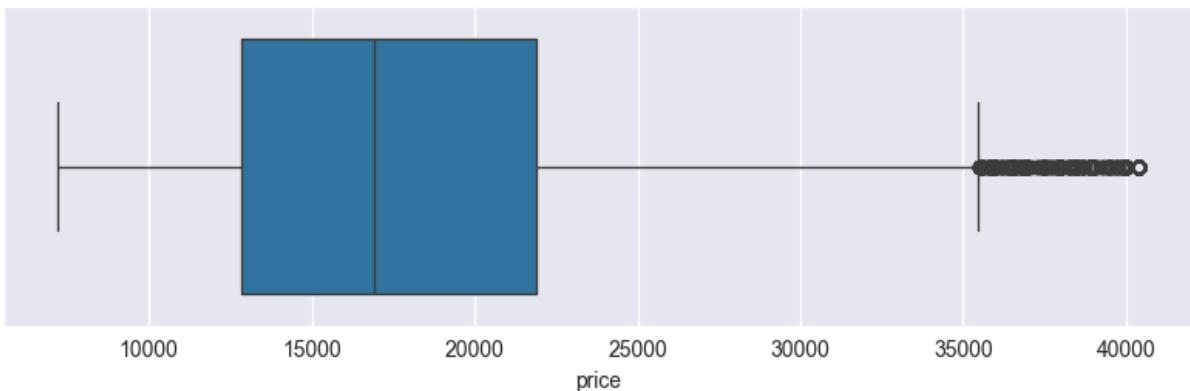
import seaborn as sns
import matplotlib.pyplot as plt

for col in numeric_cols:
    plt.figure(figsize=(8,3))
    sns.boxplot(x=df[col])
    plt.title(f"Boxplot After Outlier Handling: {col}")
    plt.tight_layout()
    plt.show()
```

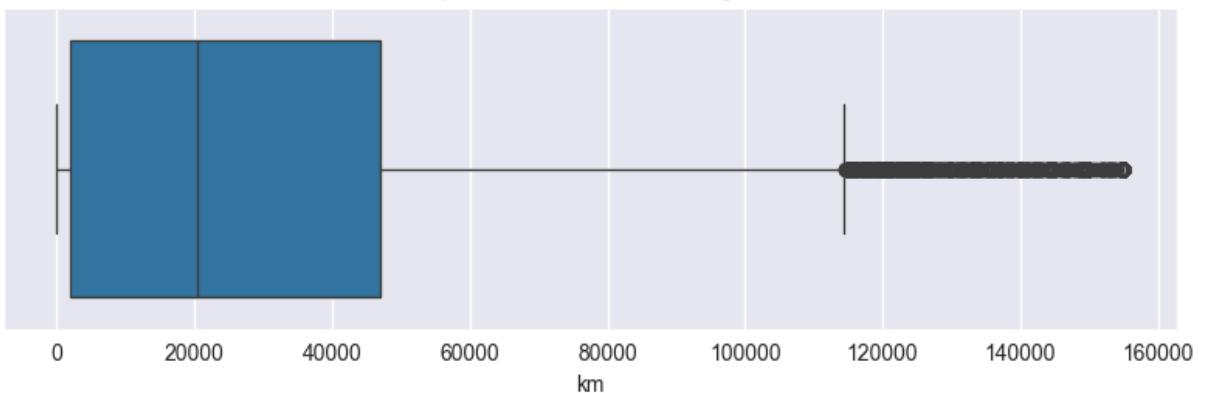
Outliers remaining after winsorization:

```
price          221
km            297
Gears          0
age            0
Previous_Owners 554
hp_kW          0
Inspection_new 0
Displacement_cc 0
Weight_kg        0
cons_comb        0
price_log        22
dtype: int64
```

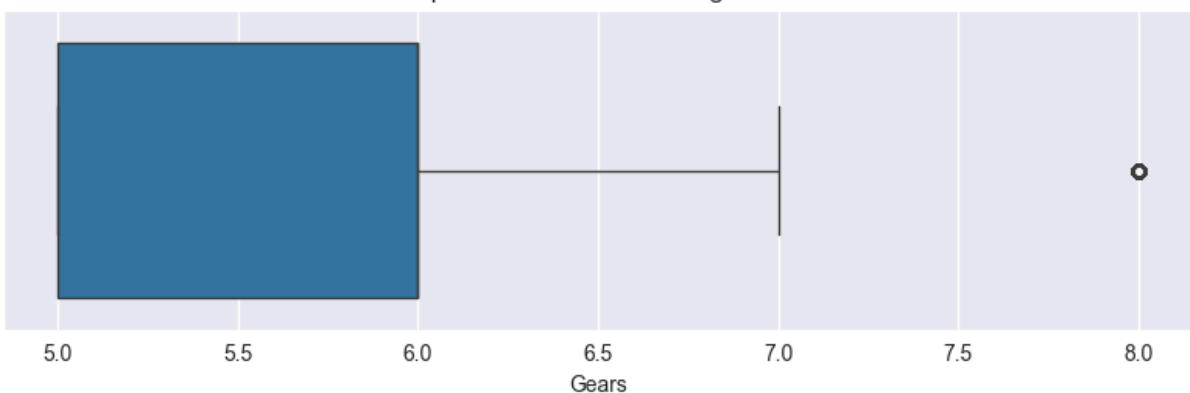
Boxplot After Outlier Handling: price



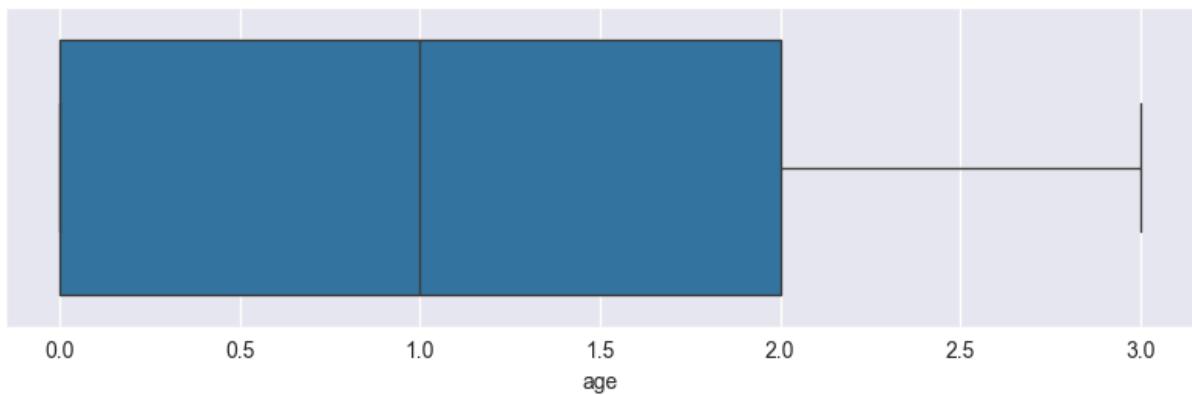
Boxplot After Outlier Handling: km



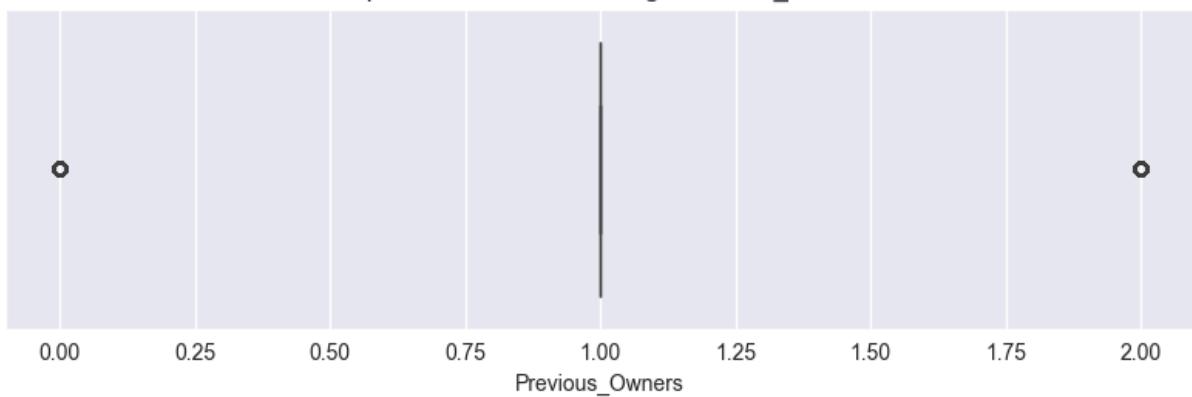
Boxplot After Outlier Handling: Gears



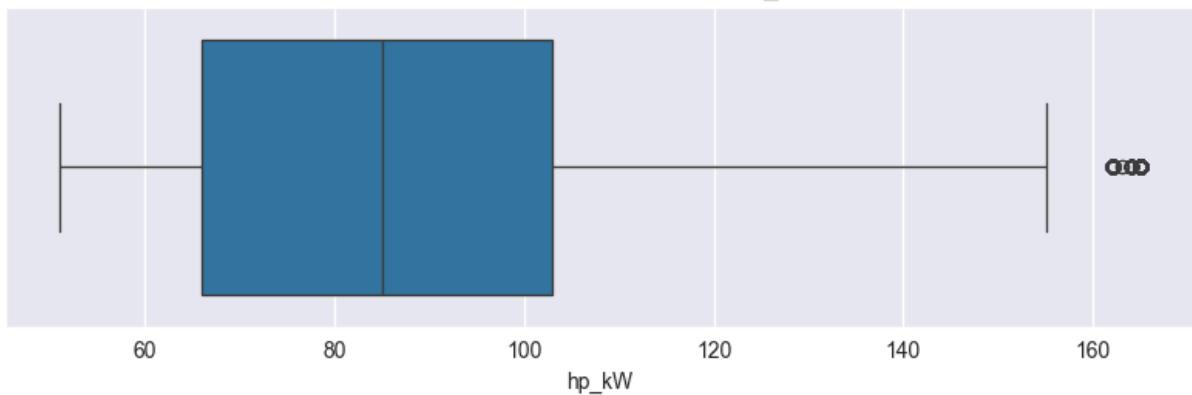
Boxplot After Outlier Handling: age



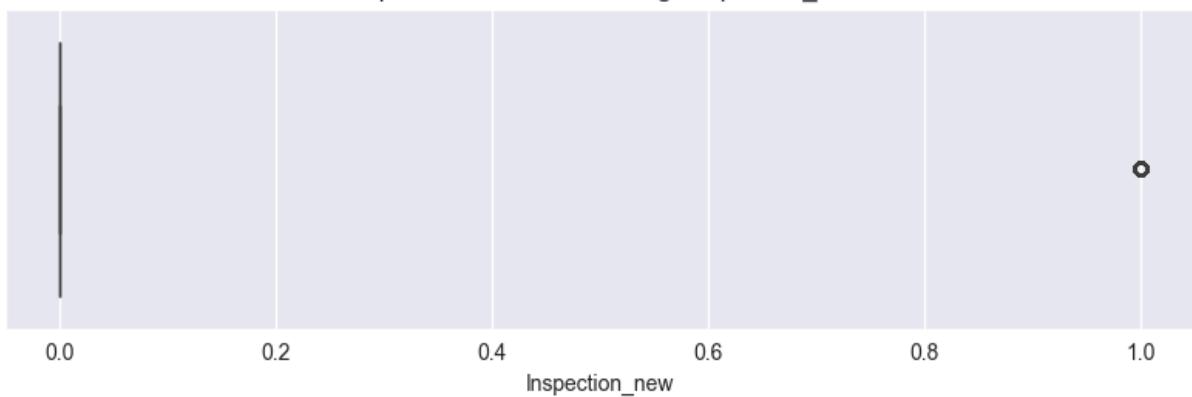
Boxplot After Outlier Handling: Previous_Owners



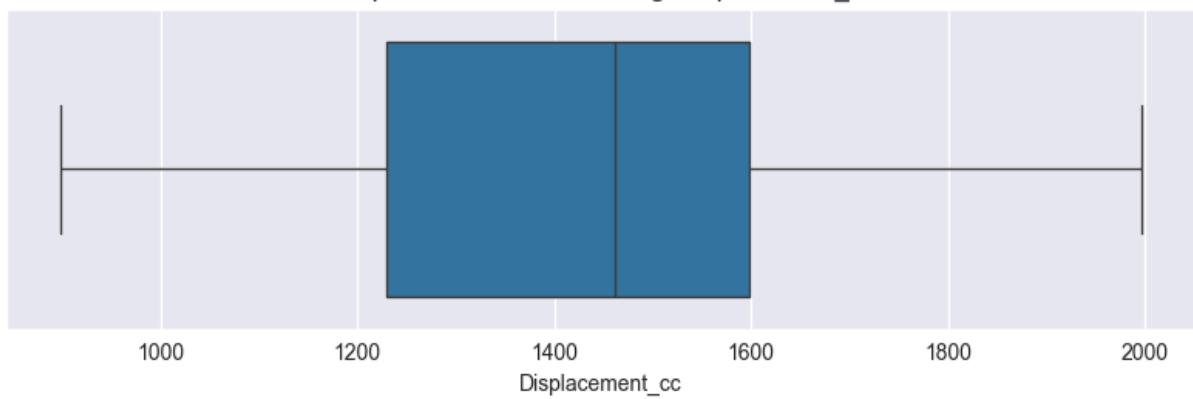
Boxplot After Outlier Handling: hp_kW



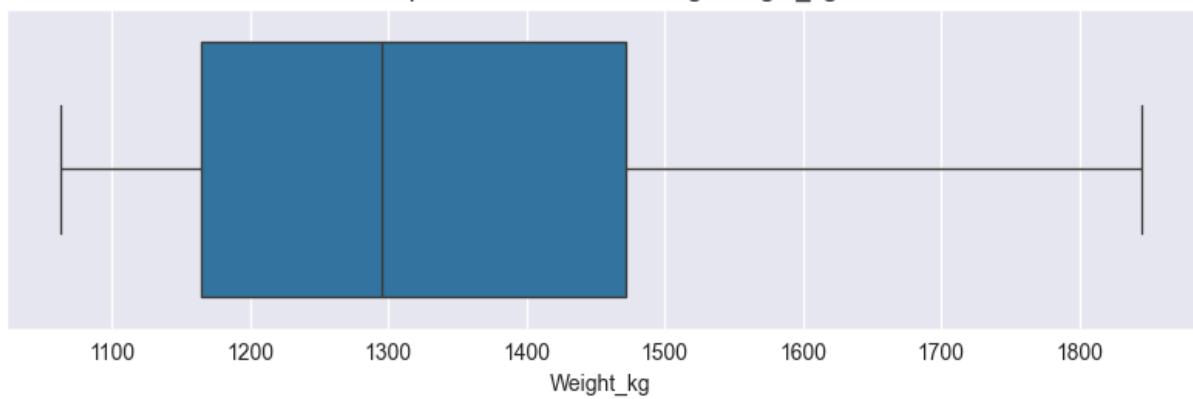
Boxplot After Outlier Handling: Inspection_new



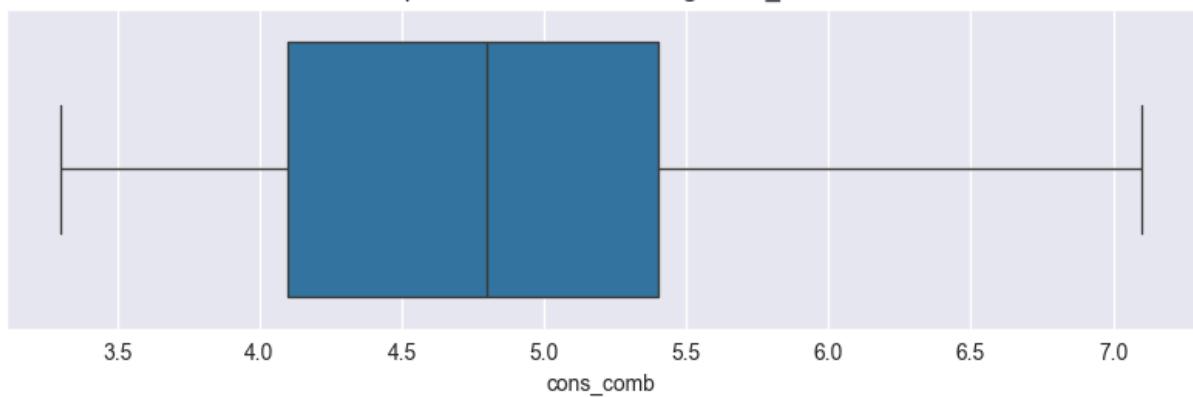
Boxplot After Outlier Handling: Displacement_cc



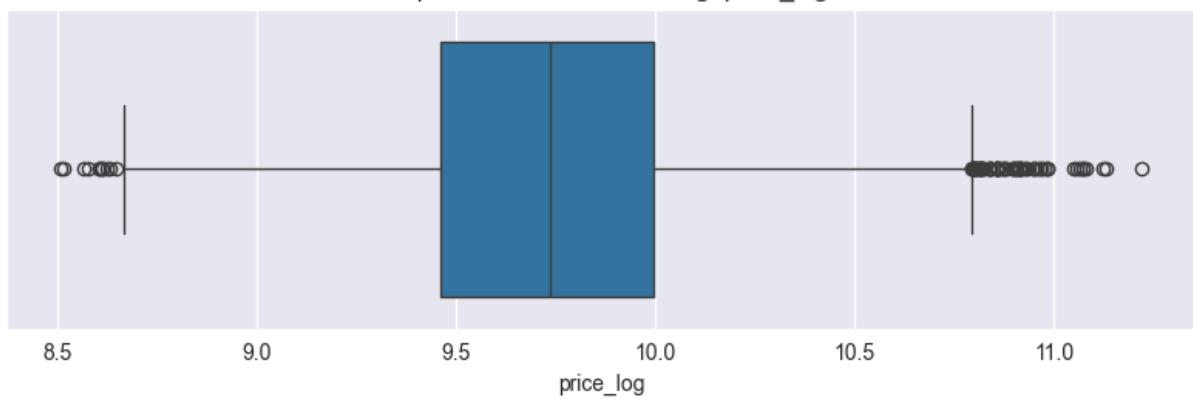
Boxplot After Outlier Handling: Weight_kg



Boxplot After Outlier Handling: cons_comb



Boxplot After Outlier Handling: price_log



2.4 Feature Engineering [11 marks]

2.4.1

Fix any redundant columns and create new ones if needed.

```
In [133]: # Fix/create columns as needed
```

2.4.2 [4 marks]

Analysis and feature engineering on `['Comfort_Convenience', 'Entertainment_Media', 'Extras', 'Safety_Security']`.

These columns contains lists of features present. Decide on how to include these features in the predictors.

```
In [134]: # Check unique values in each feature spec column
spec_cols = ['Comfort_Convenience', 'Entertainment_Media', 'Extras', 'Safety_Security']

for col in spec_cols:
    print(f"\n--- Unique sample values from {col} ---")
    print(df[col].dropna().astype(str).head(10))

def split_features(x):
    if pd.isna(x):
        return []
    return [f.strip().lower() for f in str(x).split(",") if f.strip() != ""]

for col in spec_cols:
    df[col + "_list"] = df[col].apply(split_features)

from collections import Counter

feature_counts = {}

for col in spec_cols:
    counter = Counter()
    df[col + "_list"].apply(lambda items: counter.update(items))
    feature_counts[col] = counter

for col, counter in feature_counts.items():
    print(f"\n===== Feature Frequencies in {col} =====")
    freq_df = pd.DataFrame(counter.items(), columns=["feature", "count"])
    print(freq_df.head(20))

    N = len(df)
    LOWER = 0.05 * N
    UPPER = 0.40 * N

    selected_features = []

    for col, counter in feature_counts.items():
```

```
selected = [feat for feat, cnt in counter.items() if LOWER <= cnt <= UPPER]
selected_features[col] = selected
print(f"\nSelected features from {col}:")
print(selected)# top 20

for col, feats in selected_features.items():
    for feat in feats:
        new_col = f"{col}_{feat}".replace(" ", "_")
        df[new_col] = df[col + "_list"].apply(lambda x: 1 if feat in x else
```

```
--- Unique sample values from Comfort_Convenience ---
0                               Other
1                               Other
2                               Other
3                               Other
4                               Other
5                               Other
6                               Other
7                               Other
8                               Other
9   Air conditioning,Armrest,Automatic climate con...
Name: Comfort_Convenience, dtype: object
```

```
--- Unique sample values from Entertainment_Media ---
0   Bluetooth,Hands-free equipment,On-board comput...
1                               Other
2                               Other
3   Bluetooth,CD player,Hands-free equipment,MP3,0...
4   Bluetooth,CD player,Hands-free equipment,MP3,0...
5   Bluetooth,Hands-free equipment,On-board comput...
6                               Other
7                               Other
8                               Radio
9                               Radio
Name: Entertainment_Media, dtype: object
```

```
--- Unique sample values from Extras ---
0   Alloy wheels,Catalytic Converter,Voice Control
1                               Other
2                               Alloy wheels,Voice Control
3   Alloy wheels,Sport seats,Voice Control
4                               Other
5                               Other
6                               Other
7                               Alloy wheels
8                               Alloy wheels
9                               Alloy wheels
Name: Extras, dtype: object
```

```
--- Unique sample values from Safety_Security ---
0   ABS,Central door lock,Daytime running lights,D...
1                               Other
2   ABS,Central door lock,Daytime running lights,D...
3                               Other
4                               Other
5   ABS,Central door lock,Daytime running lights,D...
6                               Other
7   ABS,Central door lock,Daytime running lights,D...
8                               Other
9   ABS,Central door lock,Daytime running lights,D...
Name: Safety_Security, dtype: object
```

```
===== Feature Frequencies in Comfort_Convenience =====
                                         feature  count
0                               other    15019
14                          power windows     896
```

1	air conditioning	896
5	electrical side mirrors	736
18	hill holder	549
3	automatic climate control	508
4	cruise control	508
9	multi-function steering wheel	508
11	park distance control	508
13	parking assist system sensors rear	348
17	start-stop system	348
16	seat heating	348
15	rain sensor	348
12	parking assist system sensors front	348
10	navigation system	348
8	lumbar support	348
7	light sensor	348
6	leather steering wheel	348
2	armrest	348
19	electrically adjustable seats	161

===== Feature Frequencies in Entertainment_Media =====

	feature	count
2	on-board computer	9146
3	radio	9124
0	bluetooth	8315
1	hands-free equipment	7027
4	other	5940
8	usb	5928
6	mp3	3771
5	cd player	2621
9	digital radio	712
7	sound system	690

===== Feature Frequencies in Extras =====

	feature	count
0	alloy wheels	10217
3	other	4710
6	touch screen	1994
2	voice control	1747
7	roof rack	1480
1	catalytic converter	930
4	sport seats	842
5	sport suspension	410

===== Feature Frequencies in Safety_Security =====

	feature	count
14	other	13120
0	abs	2795
1	central door lock	2795
2	daytime running lights	2795
3	driver-side airbag	2795
4	electronic stability control	2795
7	isofix	2795
8	passenger-side airbag	2795
9	power steering	2795
10	side airbag	2795
11	tire pressure monitoring system	2795

```

12          traction control    2795
6              immobilizer     2489
5              fog lights     1371
15      led daytime running lights   659
13          xenon headlights    621

```

Selected features from Comfort_Convenience:
['air conditioning', 'power windows']

Selected features from Entertainment_Media:
['other', 'cd player', 'mp3', 'usb']

Selected features from Extras:
['catalytic converter', 'voice control', 'other', 'sport seats', 'touch screen', 'roof rack']

Selected features from Safety_Security:
['abs', 'central door lock', 'daytime running lights', 'driver-side airbag',
'electronic stability control', 'fog lights', 'immobilizer', 'isofix', 'pass
enger-side airbag', 'power steering', 'side airbag', 'tire pressure monitori
ng system', 'traction control']

Out of these features, we will check the ones which are present in most of the cars or
are absent from most of the cars. These kinds of features can be removed as they just
increase the dimensionality without explaining the variance.

In [135...]

```
# Drop features from df
df.drop(columns=spec_cols + [col + "_list" for col in spec_cols], inplace=True)
df.head(5)
```

Out[135...]

	make_model	body_type	price	vat	km	Type	Fuel	Gear
0	Audi A1	Sedans	15770.0	VAT deductible	56013.0	Used_Regular	Diesel	7.
1	Audi A1	Sedans	14500.0	Price negotiable	80000.0	Used_Regular	Benzine	7.
2	Audi A1	Sedans	14640.0	VAT deductible	83450.0	Used_Regular	Diesel	7.
3	Audi A1	Sedans	14500.0	VAT deductible	73000.0	Used_Regular	Diesel	6.
4	Audi A1	Sedans	16790.0	VAT deductible	16200.0	Used_Regular	Diesel	7.

5 rows × 45 columns

2.4.3 [3 marks]

Perform feature encoding.

In [136..

```
import pandas as pd
import numpy as np
from sklearn.preprocessing import OneHotEncoder

all_cat_cols = df.select_dtypes(include=["object", "category"]).columns.tolist()

all_cat_cols = [c for c in all_cat_cols if c not in ["price", "price_log"]]

numeric_cols = df.select_dtypes(include=["int64", "float64"]).columns.tolist

# Remove target
numeric_cols = [c for c in numeric_cols if c not in ["price", "price_log"]]

CARDINALITY_LIMIT = 10 # <=10 → OneHot; >10 → Frequency Encoding

low_card_cols = [c for c in all_cat_cols if df[c].nunique() <= CARDINALITY_L
high_card_cols = [c for c in all_cat_cols if df[c].nunique() > CARDINALITY_L

print("Low-cardinality columns:", low_card_cols)
print("High-cardinality columns:", high_card_cols)

N = len(df)
for col in high_card_cols:
    freq = df[col].value_counts(dropna=False) / N
    df[col + "_freq"] = df[col].map(freq).fillna(0)

df.drop(columns=high_card_cols, inplace=True)

from sklearn.preprocessing import OneHotEncoder

ohe = OneHotEncoder(sparse_output=False, handle_unknown="ignore")

ohe_array = ohe.fit_transform(df[low_card_cols])
ohe_cols = ohe.get_feature_names_out(low_card_cols)

df_ohe = pd.DataFrame(ohe_array, columns=ohe_cols, index=df.index)

# Drop original low-card categorical columns
df.drop(columns=low_card_cols, inplace=True)

df_encoded = pd.concat([df, df_ohe], axis=1)

print("Final encoded dataframe shape:", df_encoded.shape)
df_encoded.head()

df_encoded = pd.concat([df, df_ohe], axis=1)
```

```

print("Final encoded dataframe shape:", df_encoded.shape)

df_encoded.head()

Low-cardinality columns: ['make_model', 'body_type', 'vat', 'Type', 'Fuel',
'Paint_Type', 'Upholstery_type', 'Gearing_Type', 'Drive_chain']
High-cardinality columns: []
Final encoded dataframe shape: (15915, 67)
Final encoded dataframe shape: (15915, 67)

```

Out[136...]

	price	km	Gears	age	Previous_Owners	hp_kW	Inspection_new	Displace
0	15770.0	56013.0	7.0	3.0		2.0	66.0	1
1	14500.0	80000.0	7.0	2.0		1.0	141.0	0
2	14640.0	83450.0	7.0	3.0		1.0	85.0	0
3	14500.0	73000.0	6.0	3.0		1.0	66.0	0
4	16790.0	16200.0	7.0	3.0		1.0	66.0	1

5 rows × 67 columns

2.4.4 [2 marks]

Split the data into training and testing sets.

In [137...]

```

# Split data
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
import pandas as pd
import numpy as np

# Ensure 'price_log' exists – create it if missing
if "price_log" not in df.columns and "price_log" not in df_encoded.columns:
    df["price_log"] = np.log1p(df["price"])    # safe log transform
    print("price_log created successfully.")

# If df_encoded exists but doesn't have price_log, add it
if "price_log" not in df_encoded.columns:
    df_encoded["price_log"] = np.log1p(df_encoded["price"])
    print("price_log added to df_encoded.")

X = df_encoded.drop(columns=["price", "price_log"], errors='ignore')
y = df_encoded["price_log"] # using the transformed target

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.20, random_state=42
)

print("Train shape:", X_train.shape)
print("Test shape:", X_test.shape)

```

```
Train shape: (12732, 65)
Test  shape: (3183, 65)
```

2.4.5 [2 marks]

Scale the features.

```
In [138]: # Scale features

numeric_cols = X_train.select_dtypes(include=['int64', 'float64']).columns.to_list()
print("Numeric columns to scale:", numeric_cols)

scaler = StandardScaler()
scaler.fit(X_train[numeric_cols])

X_train_scaled = X_train.copy()
X_test_scaled = X_test.copy()

X_train_scaled[numeric_cols] = scaler.transform(X_train[numeric_cols])
X_test_scaled[numeric_cols] = scaler.transform(X_test[numeric_cols])

print("Scaling completed successfully.")


```

Numeric columns to scale: ['km', 'Gears', 'age', 'Previous_Owners', 'hp_kw', 'Inspection_new', 'Displacement_cc', 'Weight_kg', 'cons_comb', 'Comfort_Convenience_air_conditioning', 'Comfort_Convenience_power_windows', 'Entertainment_Media_other', 'Entertainment_Media_cd_player', 'Entertainment_Media_mp3', 'Entertainment_Media_usb', 'Extras_catalytic_converter', 'Extras_voice_control', 'Extras_other', 'Extras_sport_seats', 'Extras_touch_screen', 'Extras_roof_rack', 'Safety_Security_abs', 'Safety_Security_central_door_lock', 'Safety_Security_daytime_running_lights', 'Safety_Security_driver-side_airbag', 'Safety_Security_electronic_stability_control', 'Safety_Security_fog_lights', 'Safety_Security_immobilizer', 'Safety_Security_isofix', 'Safety_Security_passenger-side_airbag', 'Safety_Security_power_steering', 'Safety_Security_side_airbag', 'Safety_Security_tire_pressure_monitoring_system', 'Safety_Security_traction_control', 'make_model_Audi A1', 'make_model_Audi A3', 'make_model_Opel Astra', 'make_model_Opel Corsa', 'make_model_Opel Insignia', 'make_model_Other', 'make_model_Renault Clio', 'make_model_Renault Espace', 'body_type_Compact', 'body_type_Other', 'body_type_Sedans', 'body_type_Station_wagon', 'body_type_Van', 'vat_Price_negotiable', 'vat_VAT_deductible', 'Type_Newly_New', 'Type_Used-Regular', 'Fuel_Benzine', 'Fuel_Diesel', 'Fuel_Other', 'Paint_Type_Metallic', 'Paint_Type_Other', 'Paint_Type_Uni/basic', 'Upholstery_type_Cloth', 'Upholstery_type_Part/Full Leather', 'Gearing_Type_Automatic', 'Gearing_Type_Manual', 'Gearing_Type_Semi-automatic', 'Drive_chain_4WD', 'Drive_chain_Other', 'Drive_chain_front']

Scaling completed successfully.

3 Linear Regression Models [35 marks]

3.1 Baseline Linear Regression Model [10 marks]

3.1.1 [5 marks]

Build and fit a basic linear regression model. Perform evaluation using suitable metrics.

```
In [139...]: # Initialise and train model
from sklearn.linear_model import LinearRegression, RidgeCV, LassoCV, ElasticNet
from sklearn.metrics import mean_squared_error, r2_score
import numpy as np

lr_model = LinearRegression()

lr_model.fit(X_train_scaled, y_train)

print("Models trained successfully.")
```

Models trained successfully.

```
In [140...]: # Evaluate the model's performance

# Predictions
pred_lr = lr_model.predict(X_test_scaled)

# RMSE & R² function
def evaluate(y_true, y_pred):
    rmse = np.sqrt(mean_squared_error(y_true, y_pred))
    r2 = r2_score(y_true, y_pred)
    return rmse, r2

# Evaluate models
results = {
    "Linear Regression": evaluate(y_test, pred_lr),
}

for model, (rmse, r2) in results.items():
    print(f"{model:20s} → RMSE: {rmse:.4f}, R²: {r2:.4f}")
```

Linear Regression → RMSE: 0.1137, R²: 0.9191

3.1.2 [5 marks]

Analyse residuals and check other assumptions of linear regression.

Check for linearity by analysing residuals vs predicted values

```
In [141...]: # Linearity check: Plot residuals vs fitted values

import matplotlib.pyplot as plt
import seaborn as sns
```

```

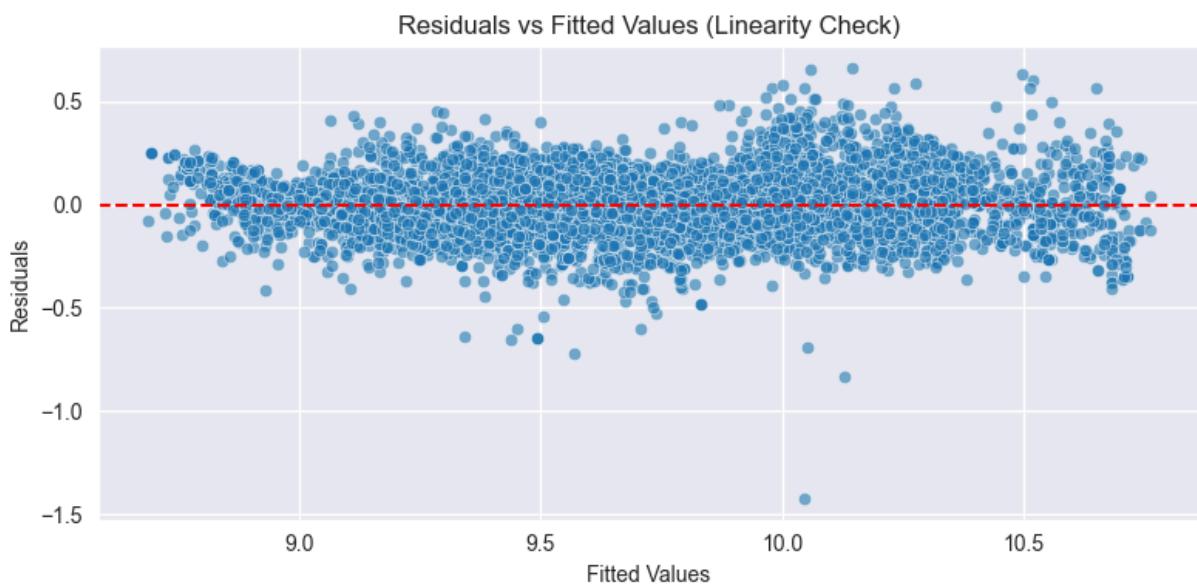
fitted_vals = lr_model.predict(X_train_scaled)
residuals = y_train - fitted_vals

plt.figure(figsize=(8, 4))
sns.scatterplot(x=fitted_vals, y=residuals, alpha=0.6)

# Add horizontal zero line
plt.axhline(0, color='red', linestyle='--')

plt.title("Residuals vs Fitted Values (Linearity Check)")
plt.xlabel("Fitted Values")
plt.ylabel("Residuals")
plt.tight_layout()
plt.show()

```



Check normality in residual distribution

```

In [142]: # Check the normality of residuals by plotting their distribution

import matplotlib.pyplot as plt
import seaborn as sns

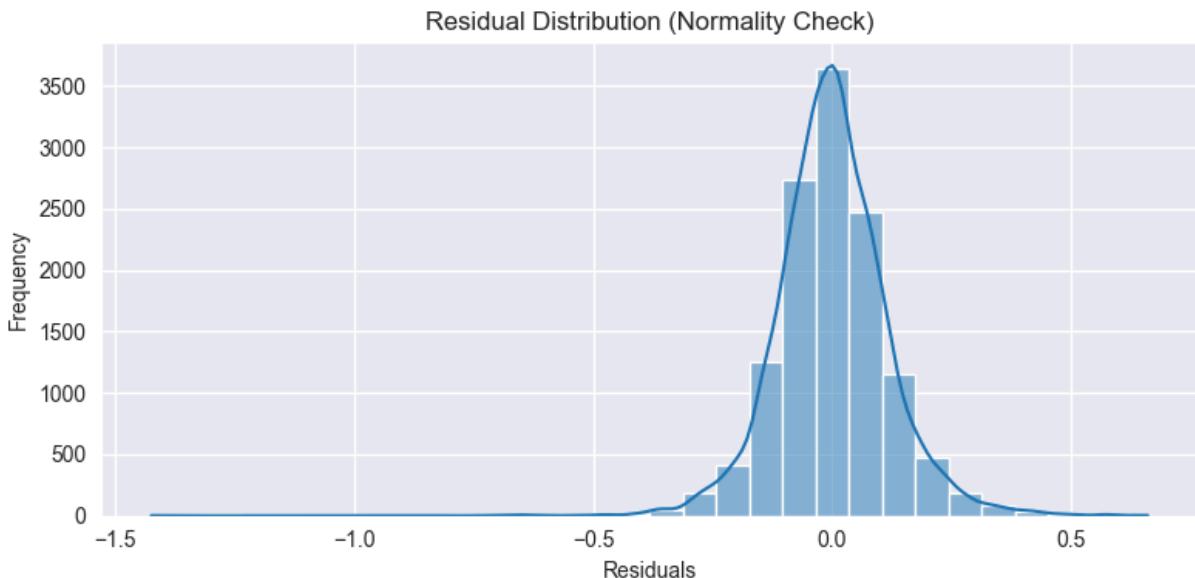
fitted_vals = lr_model.predict(X_train_scaled)
residuals = y_train - fitted_vals

plt.figure(figsize=(8, 4))
sns.histplot(residuals, bins=30, kde=True)

plt.title("Residual Distribution (Normality Check)")
plt.xlabel("Residuals")
plt.ylabel("Frequency")

```

```
plt.tight_layout()  
plt.show()
```



Check multicollinearity using Variance Inflation Factor (VIF) and handle features with high VIF.

In [143]:

```
# Check for multicollinearity and handle  
  
import pandas as pd  
import numpy as np  
from statsmodels.stats.outliers_influence import variance_inflation_factor  
  
X_vif = X_train_scaled.copy()  
  
X_vif = X_vif.replace([np.inf, -np.inf], np.nan).fillna(0)  
  
vif_data = pd.DataFrame()  
vif_data["feature"] = X_vif.columns  
vif_data["VIF"] = [variance_inflation_factor(X_vif.values, i) for i in range(X_vif.shape[1])]  
  
vif_data = vif_data.sort_values(by="VIF", ascending=False)  
vif_data.head(20)
```

/Library/Frameworks/Python.framework/Versions/3.13/lib/python3.13/site-packages/statsmodels/stats/outliers_influence.py:197: RuntimeWarning: divide by zero encountered in scalar divide
vif = 1. / (1. - r_squared_i)

Out [143...]

		feature	VIF
32	Safety_Security_tire_pressure_monitoring_system	inf	
33	Safety_Security_traction_control	inf	
35	make_model_Audi A3	inf	
36	make_model_Opel Astra	inf	
37	make_model_Opel Corsa	inf	
38	make_model_Opel Insignia	inf	
39	make_model_Other	inf	
40	make_model_Renault Clio	inf	
41	make_model_Renault Espace	inf	
42	body_type_Compact	inf	
43	body_type_Other	inf	
44	body_type_Sedans	inf	
45	body_type_Station wagon	inf	
46	body_type_Van	inf	
47	vat_Price negotiable	inf	
48	vat_VAT deductible	inf	
49	Type_Nearly_New	inf	
50	Type_Used-Regular	inf	
51	Fuel_Benzine	inf	
52	Fuel_Diesel	inf	

In [144...]

```
# Remove features with VIF > 10 (threshold can be adjusted)
high_vif_features = vif_data[vif_data["VIF"] > 10]["feature"].tolist()

print("Features to remove due to high VIF:", high_vif_features)

X_train_vif = X_train_scaled.drop(columns=high_vif_features, errors='ignore')
X_test_vif = X_test_scaled.drop(columns=high_vif_features, errors='ignore')

X_vif2 = X_train_vif.copy()
X_vif2 = X_vif2.replace([np.inf, -np.inf], np.nan).fillna(0)

vif_data2 = pd.DataFrame()
vif_data2["feature"] = X_vif2.columns
vif_data2["VIF"] = [variance_inflation_factor(X_vif2.values, i) for i in range(len(X_vif2.columns))]
vif_data2.sort_values(by="VIF", ascending=False).head(20)
```

Features to remove due to high VIF: ['Safety_Security_tire_pressure_monitoring_system', 'Safety_Security_traction_control', 'make_model_Audi A3', 'make_model_Opel Astra', 'make_model_Opel Corsa', 'make_model_Opel Insignia', 'make_model_Other', 'make_model_Renault Clio', 'make_model_Renault Espace', 'body_type_Compact', 'body_type_Other', 'body_type_Sedans', 'body_type_Station wagon', 'body_type_Van', 'vat_Price negotiable', 'vat_VAT deductible', 'Type_Nearly_New', 'Type_Used-Regular', 'Fuel_Benzine', 'Fuel_Diesel', 'Fuel_Other', 'Paint_Type_Metallic', 'Paint_Type_Other', 'Paint_Type_Uni/basic', 'Upholstery_type_Cloth', 'Upholstery_type_Part/Full Leather', 'Gearing_Type_Automatic', 'Gearing_Type_Manual', 'Gearing_Type_Semi-automatic', 'Drive_chain_4WD', 'Drive_chain_Other', 'make_model_Audi A1', 'Drive_chain_front', 'Safety_Security_daytime_running_lights', 'Safety_Security_electronic_stability_control', 'Safety_Security_side_airbag', 'Safety_Security_abs', 'Safety_Security_central_door_lock', 'Comfort_Convenience_power_windows', 'Safety_Security_driver-side_airbag', 'Comfort_Convenience_air_conditioning', 'Safety_Security_isofix', 'Safety_Security_passenger-side_airbag', 'Safety_Security_power_steering']

Out[144...]

	feature	VIF
4	hp_kw	3.168558
2	age	2.873470
0	km	2.714429
20	Safety_Security_immobilizer	2.233195
7	Weight_kg	2.174262
19	Safety_Security_fog_lights	2.108772
12	Entertainment_Media_usb	2.086551
6	Displacement_cc	2.018425
11	Entertainment_Media_mp3	1.759085
9	Entertainment_Media_other	1.746675
8	cons_comb	1.599391
10	Entertainment_Media_cd_player	1.548155
15	Extras_other	1.516043
1	Gears	1.371631
17	Extras_touch_screen	1.325677
14	Extras_voice_control	1.221952
3	Previous_Owners	1.153945
5	Inspection_new	1.144758
16	Extras_sport_seats	1.139439
18	Extras_roof_rack	1.139352

3.2 Ridge Regression Implementation [10 marks]

3.2.1 [2 marks]

Define a list of random alpha values

```
In [168...]: # List of alphas to tune for Ridge regularisation
import numpy as np

ridge_alphas = np.logspace(-4, 4, 50)
ridge_alphas
```

```
Out[168...]: array([1.00000000e-04, 1.45634848e-04, 2.12095089e-04, 3.08884360e-04,
       4.49843267e-04, 6.55128557e-04, 9.54095476e-04, 1.38949549e-03,
       2.02358965e-03, 2.94705170e-03, 4.29193426e-03, 6.25055193e-03,
       9.10298178e-03, 1.32571137e-02, 1.93069773e-02, 2.81176870e-02,
       4.09491506e-02, 5.96362332e-02, 8.68511374e-02, 1.26485522e-01,
       1.84206997e-01, 2.68269580e-01, 3.90693994e-01, 5.68986603e-01,
       8.28642773e-01, 1.20679264e+00, 1.75751062e+00, 2.55954792e+00,
       3.72759372e+00, 5.42867544e+00, 7.90604321e+00, 1.15139540e+01,
       1.67683294e+01, 2.44205309e+01, 3.55648031e+01, 5.17947468e+01,
       7.54312006e+01, 1.09854114e+02, 1.59985872e+02, 2.32995181e+02,
       3.39322177e+02, 4.94171336e+02, 7.19685673e+02, 1.04811313e+03,
       1.52641797e+03, 2.22299648e+03, 3.23745754e+03, 4.71486636e+03,
       6.86648845e+03, 1.00000000e+04])
```

3.2.2 [4 marks]

Apply Ridge Regularisation and find the best value of alpha from the list

```
In [169...]: # Applying Ridge regression
from sklearn.linear_model import RidgeCV
from sklearn.metrics import mean_squared_error, r2_score
import numpy as np

ridge_alphas = np.logspace(-4, 4, 50)

ridge_model = RidgeCV(alphas=ridge_alphas, cv=5)

ridge_model.fit(X_train_scaled, y_train)

print("Ridge model trained successfully.")
print("Best alpha selected:", ridge_model.alpha_)

ridge_pred = ridge_model.predict(X_test_scaled)

ridge_rmse = np.sqrt(mean_squared_error(y_test, ridge_pred))
ridge_r2 = r2_score(y_test, ridge_pred)
```

```
print(f"Ridge RMSE: {ridge_rmse:.4f}")
print(f"Ridge R2: {ridge_r2:.4f}")
```

```
Ridge model trained successfully.
Best alpha selected: 7.9060432109076855
Ridge RMSE: 0.1137
Ridge R2: 0.9191
```

```
In [170]: # Plot train and test scores against alpha
import numpy as np
import matplotlib.pyplot as plt
from sklearn.linear_model import Ridge
from sklearn.metrics import r2_score

alphas = np.logspace(-4, 4, 50)

train_scores = []
test_scores = []

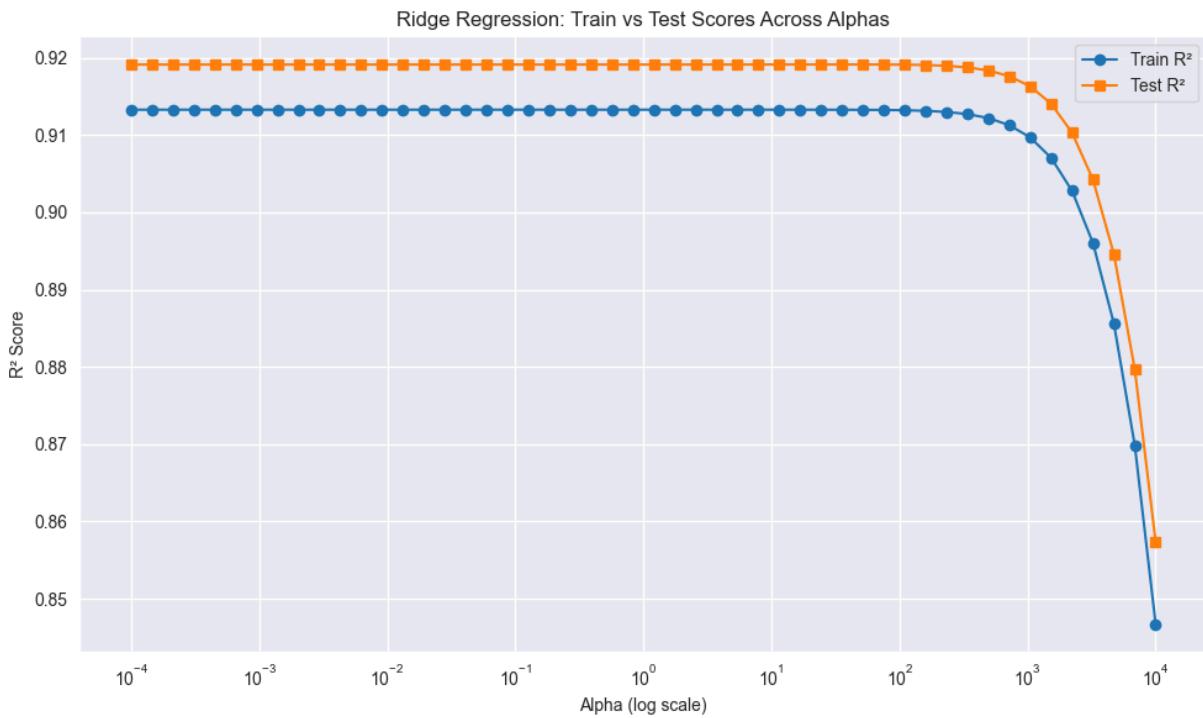
for alpha in alphas:
    ridge = Ridge(alpha=alpha)
    ridge.fit(X_train_scaled, y_train)

    train_scores.append(r2_score(y_train, ridge.predict(X_train_scaled)))
    test_scores.append(r2_score(y_test, ridge.predict(X_test_scaled)))

plt.figure(figsize=(10, 6))

plt.plot(alphas, train_scores, label="Train R2", marker='o')
plt.plot(alphas, test_scores, label="Test R2", marker='s')

plt.xscale("log")
plt.xlabel("Alpha (log scale)")
plt.ylabel("R2 Score")
plt.title("Ridge Regression: Train vs Test Scores Across Alphas")
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.show()
```



Find the best alpha value.

```
In [171]: # Best alpha value

# Best score (negative MAE)

from sklearn.linear_model import RidgeCV
from sklearn.metrics import mean_absolute_error
import numpy as np

alphas = np.logspace(-4, 4, 50)

ridge_mae_model = RidgeCV(alphas=alphas, scoring='neg_mean_absolute_error',
                          cv=5)

# Train model
ridge_mae_model.fit(X_train_scaled, y_train)

best_alpha = ridge_mae_model.alpha_
print("Best Alpha (MAE scoring):", best_alpha)

best_score_neg_mae = ridge_mae_model.best_score_
print("Best CV Score (negative MAE):", best_score_neg_mae)

# Convert negative MAE → actual MAE
best_mae = -best_score_neg_mae
print("Best CV MAE:", best_mae)
```

```
Best Alpha (MAE scoring): 35.564803062231285
Best CV Score (negative MAE): -0.08690307331877282
Best CV MAE: 0.08690307331877282
```

We will get some best value of alpha above. This however is not the most accurate value but the best value from the given list. Now we have a rough estimate of the range that best alpha falls in. Let us do another iteration over the values in a smaller range.

3.2.3 [4 marks]

Fine tune by taking a closer range of alpha based on the previous result.

```
In [176...]: # Take a smaller range of alpha to test

import numpy as np

# Use a smaller log-spaced range around typical optimal ridge α values
alphas_small = np.logspace(-2, 1, 20)  # 0.01 → 10 (much tighter)
alphas_small
```



```
Out[176...]: array([ 0.01      ,  0.0143845 ,  0.02069138,  0.02976351,  0.04281332,
       0.06158482,  0.08858668,  0.1274275 ,  0.18329807,  0.26366509,
       0.37926902,  0.54555948,  0.78475997,  1.12883789,  1.62377674,
       2.33572147,  3.35981829,  4.83293024,  6.95192796, 10.        ])
```

```
In [177...]: # Applying Ridge regression

from sklearn.linear_model import RidgeCV

ridge_mae_refined = RidgeCV(
    alphas=alphas_small,
    scoring='neg_mean_absolute_error',
    cv=5
)

ridge_mae_refined.fit(X_train_scaled, y_train)

print("Best alpha (refined search):", ridge_mae_refined.alpha_)
print("Best score (negative MAE):", ridge_mae_refined.best_score_)
print("Best MAE:", -ridge_mae_refined.best_score_)
```

```
Best alpha (refined search): 10.0
Best score (negative MAE): -0.08690678416587028
Best MAE: 0.08690678416587028
```

Plot the error-alpha graph again and find the actual optimal value for alpha.

```
In [178...]: # Plot train and test scores against alpha

import numpy as np
import matplotlib.pyplot as plt
from sklearn.linear_model import Ridge
```

```

from sklearn.linear_model import RidgeCV
from sklearn.metrics import r2_score, mean_absolute_error

alphas_small = np.logspace(-2, 1, 20)    # 0.01 → 10 (tighter range)

train_scores = []
test_scores = []

for alpha in alphas_small:
    ridge = Ridge(alpha=alpha)
    ridge.fit(X_train_scaled, y_train)

    train_scores.append(r2_score(y_train, ridge.predict(X_train_scaled)))
    test_scores.append(r2_score(y_test, ridge.predict(X_test_scaled)))

plt.figure(figsize=(10, 6))
plt.plot(alphas_small, train_scores, marker='o', label="Train R²")
plt.plot(alphas_small, test_scores, marker='s', label="Test R²")

plt.xscale("log")
plt.xlabel("Alpha (log scale)")
plt.ylabel("R² Score")
plt.title("Ridge Regression: Train & Test Scores vs Alpha (Refined Search)")
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.show()

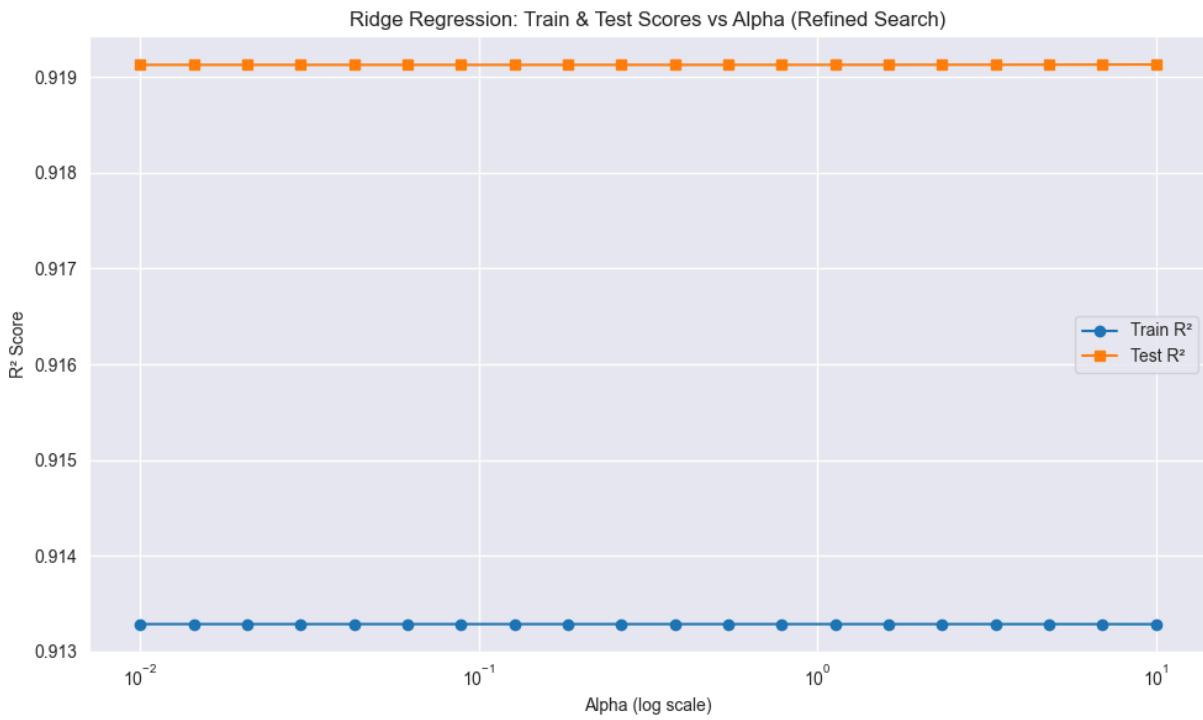
ridge_mae_refined = RidgeCV(
    alphas=alphas_small,
    scoring='neg_mean_absolute_error',
    cv=5
)

ridge_mae_refined.fit(X_train_scaled, y_train)

best_alpha = ridge_mae_refined.alpha_
best_neg_mae = ridge_mae_refined.best_score_

print("\n=====")
print(" BEST RESULTS FROM REFINED SEARCH")
print("=====")
print("Best alpha value:", best_alpha)
print("Best score (negative MAE):", best_neg_mae)
print("Best MAE:", -best_neg_mae)

```



=====

BEST RESULTS FROM REFINED SEARCH

=====

Best alpha value: 10.0
 Best score (negative MAE): -0.08690678416587028
 Best MAE: 0.08690678416587028

```
In [179]: # Set best alpha for Ridge regression
# Fit the Ridge model to get the coefficients of the fitted model

from sklearn.linear_model import Ridge

best_alpha = 10.0 # or ridge_mae_refined.alpha_
ridge_final = Ridge(alpha=best_alpha)

ridge_final.fit(X_train_scaled, y_train)

print("Final Ridge model fitted with alpha =", best_alpha)
```

Final Ridge model fitted with alpha = 10.0

```
In [180]: # Show the coefficients for each feature

ridge_coeffs = ridge_final.coef_

# Combine with feature names
coeff_df = pd.DataFrame({
    "Feature": X_train_scaled.columns,
    "Coefficient": ridge_coeffs
}).sort_values(by="Coefficient", ascending=False)
```

```
print("Top coefficients:")
coeff_df.head(15)
```

Top coefficients:

Out[180...]

	Feature	Coefficient
4	hp_kw	0.122459
35	make_model_Audi A3	0.083487
41	make_model_Renault Espace	0.069814
34	make_model_Audi A1	0.046989
61	Gearing_Type_Semi-automatic	0.029695
59	Gearing_Type_Automatic	0.027389
38	make_model_Opel Insignia	0.027108
1	Gears	0.017510
17	Extras_other	0.010877
52	Fuel_Diesel	0.009528
8	cons_comb	0.008735
58	Upholstery_type_Part/Full Leather	0.007992
13	Entertainment_Media_mp3	0.007465
19	Extras_touch_screen	0.007380
11	Entertainment_Media_other	0.006713

In [181...]

```
# Evaluate the Ridge model on the test data
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
import numpy as np

y_pred_ridge = ridge_final.predict(X_test_scaled)

ridge_rmse = np.sqrt(mean_squared_error(y_test, y_pred_ridge))
ridge_mae = mean_absolute_error(y_test, y_pred_ridge)
ridge_r2 = r2_score(y_test, y_pred_ridge)

print("==== Ridge Model Evaluation on Test Data ====")
print(f"RMSE: {ridge_rmse:.4f}")
print(f"MAE : {ridge_mae:.4f}")
print(f"R²  : {ridge_r2:.4f}")
```

==== Ridge Model Evaluation on Test Data ====

RMSE: 0.1137

MAE : 0.0855

R² : 0.9191

3.3 Lasso Regression Implementation [10 marks]

3.3.1 [2 marks]

Define a list of random alpha values

```
In [155...]: # List of alphas to tune for Lasso regularisation  
  
import numpy as np  
  
lasso_alphas = np.logspace(-4, 1, 50) # 0.0001 → 10  
lasso_alphas
```

3.3.2 [4 marks]

Apply Ridge Regularisation and find the best value of alpha from the list

```
In [182...]: # Initialise Lasso regression model  
from sklearn.linear_model import LassoCV  
import numpy as np  
  
lasso_alphas = np.logspace(-4, 1, 50) # 0.0001 → 10  
  
lasso_model = LassoCV(  
    alphas=lasso_alphas,  
    cv=5,  
    max_iter=5000,  
    random_state=42  
)  
  
print("Lasso model initialized.")
```

Lasso model initialized.

```
In [183...]: # Plot train and test scores against alpha  
  
import numpy as np  
import matplotlib.pyplot as plt  
from sklearn.linear_model import Lasso  
from sklearn.metrics import r2_score  
  
lasso_alphas = np.logspace(-4, 1, 50) # 0.0001 → 10  
  
train_scores = []  
test_scores = []  
  
for alpha in lasso_alphas:  
    lasso = Lasso(alpha=alpha, max_iter=5000, random_state=42)  
    lasso.fit(X_train_scaled, y_train)
```

```

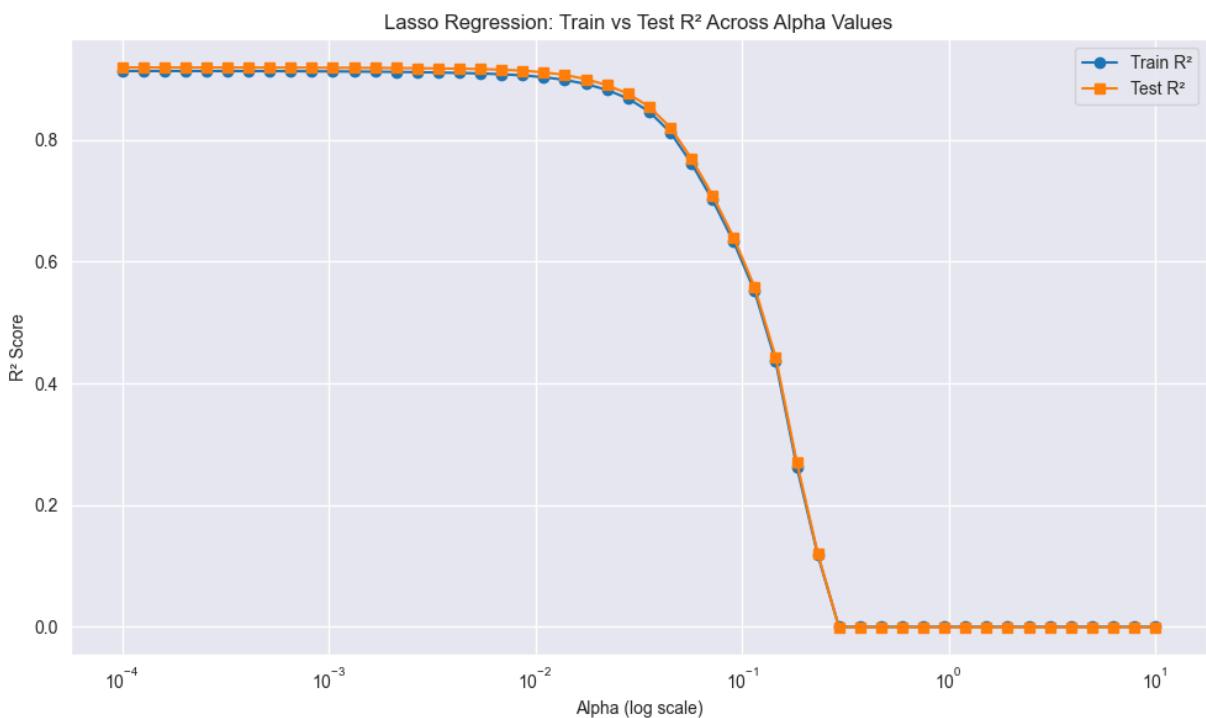
train_scores.append(r2_score(y_train, lasso.predict(X_train_scaled)))
test_scores.append(r2_score(y_test, lasso.predict(X_test_scaled)))

plt.figure(figsize=(10, 6))

plt.plot(lasso_alphas, train_scores, label="Train R2", marker="o")
plt.plot(lasso_alphas, test_scores, label="Test R2", marker="s")

plt.xscale("log")
plt.xlabel("Alpha (log scale)")
plt.ylabel("R2 Score")
plt.title("Lasso Regression: Train vs Test R2 Across Alpha Values")
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.show()

```



In [184...]:

```

# Best alpha value

# Best score (negative MAE)

from sklearn.linear_model import LassoCV
import numpy as np

lasso_alphas = np.logspace(-4, 1, 50)    # 0.0001 → 10

lasso_cv = LassoCV(
    alphas=lasso_alphas,
    cv=5,
)

```

```

        max_iter=5000,
        random_state=42,
    )

lasso_cv.fit(X_train_scaled, y_train)

best_alpha = lasso_cv.alpha_
best_neg_mae = lasso_cv.mse_path_.mean(axis=1)[list(lasso_alphas).index(best_alpha)]
print("Best alpha (Lasso):", best_alpha)
print("Best score (negative MAE):", best_neg_mae)
print("Best MAE:", -best_neg_mae)

```

Best alpha (Lasso): 0.00012648552168552957
 Best score (negative MAE): 0.15825313379090683
 Best MAE: -0.15825313379090683

3.3.3 [4 marks]

Fine tune by taking a closer range of alpha based on the previous result.

In [185... # List of alphas to tune for Lasso regularization

```

import numpy as np

# Assume best_alpha_prev is already known
best_alpha_prev = lasso_cv.alpha_

# Build a tighter window around the previous best alpha
alphas_refined = np.logspace(
    np.log10(best_alpha_prev) - 1,
    np.log10(best_alpha_prev) + 1,
    30
)

print("Refined alpha range:")
print(alphas_refined)

```

Refined alpha range:

```
[1.26485522e-05 1.48253971e-05 1.73768820e-05 2.03674833e-05
 2.38727739e-05 2.79813332e-05 3.27969849e-05 3.84414213e-05
 4.50572783e-05 5.28117394e-05 6.19007611e-05 7.25540243e-05
 8.50407385e-05 9.96764450e-05 1.16830990e-04 1.36937872e-04
 1.60505194e-04 1.88128507e-04 2.20505856e-04 2.58455420e-04
 3.02936191e-04 3.55072206e-04 4.16180948e-04 4.87806646e-04
 5.71759290e-04 6.70160378e-04 7.85496519e-04 9.20682274e-04
 1.07913381e-03 1.26485522e-03]
```

In [186... # Tuning Lasso hyperparameters

Refined Best Alpha: 0.00011683099023706688

In [187... # Plot train and test scores against alpha

```

import numpy as np
import matplotlib.pyplot as plt

```

```
from sklearn.linear_model import Lasso
from sklearn.metrics import r2_score

train_scores_refined = []
test_scores_refined = []

for alpha in alphas_refined:
    lasso_temp = Lasso(alpha=alpha, max_iter=5000, random_state=42)
    lasso_temp.fit(X_train_scaled, y_train)

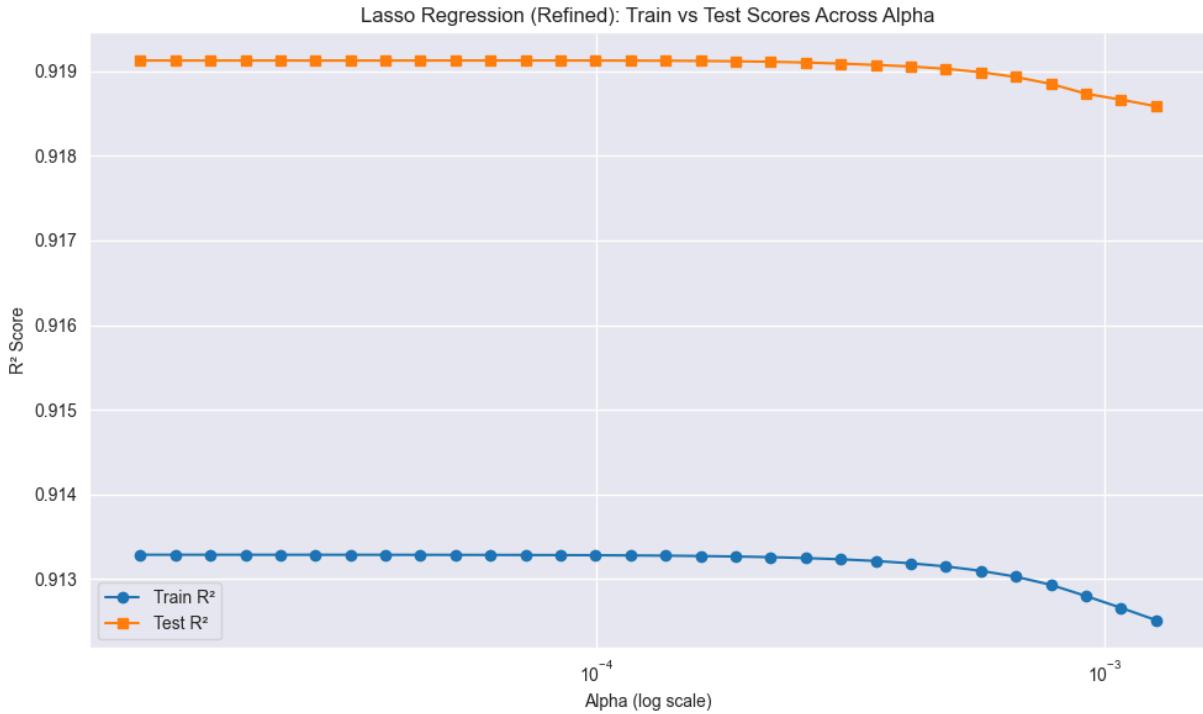
    train_scores_refined.append(r2_score(y_train, lasso_temp.predict(X_train_scaled)))
    test_scores_refined.append(r2_score(y_test, lasso_temp.predict(X_test_scaled)))

plt.figure(figsize=(10, 6))

plt.plot(alphas_refined, train_scores_refined, marker="o", label="Train R²")
plt.plot(alphas_refined, test_scores_refined, marker="s", label="Test R²")

plt.xscale("log")
plt.xlabel("Alpha (log scale)")
plt.ylabel("R² Score")
plt.title("Lasso Regression (Refined): Train vs Test Scores Across Alpha")
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.show()
```

```
/Library/Frameworks/Python.framework/Versions/3.13/lib/python3.13/site-pac
ges/sklearn/linear_model/_coordinate_descent.py:695: ConvergenceWarning: Obj
ective did not converge. You might want to increase the number of iteration
s, check the scale of the features or consider increasing regularisation. Du
ality gap: 1.574e+00, tolerance: 2.015e-01
    model = cd_fast.enet_coordinate_descent(
/Library/Frameworks/Python.framework/Versions/3.13/lib/python3.13/site-pac
ges/sklearn/linear_model/_coordinate_descent.py:695: ConvergenceWarning: Obj
ective did not converge. You might want to increase the number of iteration
s, check the scale of the features or consider increasing regularisation. Du
ality gap: 1.578e+00, tolerance: 2.015e-01
    model = cd_fast.enet_coordinate_descent(
/Library/Frameworks/Python.framework/Versions/3.13/lib/python3.13/site-pac
ges/sklearn/linear_model/_coordinate_descent.py:695: ConvergenceWarning: Obj
ective did not converge. You might want to increase the number of iteration
s, check the scale of the features or consider increasing regularisation. Du
ality gap: 1.582e+00, tolerance: 2.015e-01
    model = cd_fast.enet_coordinate_descent(
/Library/Frameworks/Python.framework/Versions/3.13/lib/python3.13/site-pac
ges/sklearn/linear_model/_coordinate_descent.py:695: ConvergenceWarning: Obj
ective did not converge. You might want to increase the number of iteration
s, check the scale of the features or consider increasing regularisation. Du
ality gap: 1.587e+00, tolerance: 2.015e-01
    model = cd_fast.enet_coordinate_descent(
/Library/Frameworks/Python.framework/Versions/3.13/lib/python3.13/site-pac
ges/sklearn/linear_model/_coordinate_descent.py:695: ConvergenceWarning: Obj
ective did not converge. You might want to increase the number of iteration
s, check the scale of the features or consider increasing regularisation. Du
ality gap: 1.592e+00, tolerance: 2.015e-01
    model = cd_fast.enet_coordinate_descent(
/Library/Frameworks/Python.framework/Versions/3.13/lib/python3.13/site-pac
ges/sklearn/linear_model/_coordinate_descent.py:695: ConvergenceWarning: Obj
ective did not converge. You might want to increase the number of iteration
s, check the scale of the features or consider increasing regularisation. Du
ality gap: 1.598e+00, tolerance: 2.015e-01
    model = cd_fast.enet_coordinate_descent(
/Library/Frameworks/Python.framework/Versions/3.13/lib/python3.13/site-pac
ges/sklearn/linear_model/_coordinate_descent.py:695: ConvergenceWarning: Obj
ective did not converge. You might want to increase the number of iteration
s, check the scale of the features or consider increasing regularisation. Du
ality gap: 1.603e+00, tolerance: 2.015e-01
    model = cd_fast.enet_coordinate_descent(
```



```
In [188]: # Best alpha value
```

```
# Best score (negative MAE)
from sklearn.linear_model import LassoCV

lasso_refined = LassoCV(
    alphas=alphas_refined,
    cv=5,
    max_iter=5000,
    random_state=42
)

lasso_refined.fit(X_train_scaled, y_train)

best_alpha_refined = lasso_refined.alpha_
print("Refined Best Alpha:", best_alpha_refined)
```

Reffined Best Alpha: 0.00011683099023706688

```
In [189]: # Set best alpha for Lasso regression
```

```
# Fit the Lasso model on scaled training data
# Get the coefficients of the fitted model
from sklearn.linear_model import Lasso
import pandas as pd

best_alpha_lasso = lasso_refined.alpha_
print("Using Best Alpha for Final Lasso:", best_alpha_lasso)

lasso_final = Lasso(alpha=best_alpha_lasso, max_iter=5000, random_state=42)
```

```

lasso_final.fit(X_train_scaled, y_train)

print("Final Lasso model fitted successfully.")

lasso_coeffs = lasso_final.coef_

coeff_df = pd.DataFrame({
    "Feature": X_train_scaled.columns,
    "Coefficient": lasso_coeffs
}).sort_values(by="Coefficient", ascending=False)

print("\nTop Lasso Coefficients (Positive Impact):")
print(coeff_df.head(10))

print("\nTop Lasso Coefficients (Negative Impact):")
print(coeff_df.tail(10))

```

Using Best Alpha for Final Lasso: 0.00011683099023706688
Final Lasso model fitted successfully.

Top Lasso Coefficients (Positive Impact):

	Feature	Coefficient
4	hp_kw	0.121955
35	make_model_Audi A3	0.054635
41	make_model_Renault Espace	0.052529
61	Gearing_Type_Semi-automatic	0.020237
34	make_model_Audi A1	0.020218
1	Gears	0.017365
17	Extras_other	0.010524
52	Fuel_Diesel	0.009052
49	Type_Nearly_New	0.008773
8	cons_comb	0.007948

Top Lasso Coefficients (Negative Impact):

	Feature	Coefficient
51	Fuel_Benzine	-0.008516
6	Displacement_cc	-0.010708
57	Upholstery_type_Cloth	-0.015932
39	make_model_Other	-0.018404
36	make_model_Opel Astra	-0.059091
60	Gearing_Type_Manual	-0.065041
0	km	-0.088445
2	age	-0.109603
40	make_model_Renault Clio	-0.110660
37	make_model_Opel Corsa	-0.132418

In [190]: # Check the coefficients for each feature

```

lasso_coeff_df = pd.DataFrame({
    "Feature": X_train_scaled.columns,
    "Coefficient": lasso_final.coef_
})

```

```

lasso_coeff_df_sorted = lasso_coeff_df.sort_values(
    by="Coefficient",
    ascending=False
).reset_index(drop=True)

print("===== LASSO COEFFICIENTS (SORTED) =====")
display(lasso_coeff_df_sorted)

```

===== LASSO COEFFICIENTS (SORTED) =====

	Feature	Coefficient
0	hp_kw	0.121955
1	make_model_Audi A3	0.054635
2	make_model_Renault Espace	0.052529
3	Gearing_Type_Semi-automatic	0.020237
4	make_model_Audi A1	0.020218
...
60	Gearing_Type_Manual	-0.065041
61	km	-0.088445
62	age	-0.109603
63	make_model_Renault Clio	-0.110660
64	make_model_Opel Corsa	-0.132418

65 rows × 2 columns

```

In [191]: # Evaluate the Lasso model on the test data

from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
import numpy as np

y_pred_lasso = lasso_final.predict(X_test_scaled)

lasso_rmse = np.sqrt(mean_squared_error(y_test, y_pred_lasso))
lasso_mae = mean_absolute_error(y_test, y_pred_lasso)
lasso_r2 = r2_score(y_test, y_pred_lasso)

print("===== LASSO MODEL EVALUATION (TEST SET) =====")
print(f"RMSE : {lasso_rmse:.4f}")
print(f"MAE : {lasso_mae:.4f}")
print(f"R² : {lasso_r2:.4f}")

```

===== LASSO MODEL EVALUATION (TEST SET) =====

RMSE : 0.1137
MAE : 0.0856
R² : 0.9191

3.4 Regularisation Comparison & Analysis [5 marks]

3.4.1 [2 marks]

Compare the evaluation metrics for each model.

```
In [192...]: # Compare metrics for each model
import pandas as pd
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
import numpy as np

y_pred_lr      = lr_model.predict(X_test_scaled)
y_pred_ridge   = ridge_final.predict(X_test_scaled)
y_pred_lasso   = lasso_final.predict(X_test_scaled)

metrics = {
    "Model": ["Linear Regression", "Ridge Regression", "Lasso Regression"],
    "RMSE": [
        np.sqrt(mean_squared_error(y_test, y_pred_lr)),
        np.sqrt(mean_squared_error(y_test, y_pred_ridge)),
        np.sqrt(mean_squared_error(y_test, y_pred_lasso))
    ],
    "MAE": [
        mean_absolute_error(y_test, y_pred_lr),
        mean_absolute_error(y_test, y_pred_ridge),
        mean_absolute_error(y_test, y_pred_lasso)
    ],
    "R² Score": [
        r2_score(y_test, y_pred_lr),
        r2_score(y_test, y_pred_ridge),
        r2_score(y_test, y_pred_lasso)
    ]
}

metrics_df = pd.DataFrame(metrics)
metrics_df
```

```
Out[192...]:
```

	Model	RMSE	MAE	R ² Score
0	Linear Regression	0.113697	0.085551	0.919127
1	Ridge Regression	0.113694	0.085550	0.919130
2	Lasso Regression	0.113697	0.085562	0.919126

3.4.2 [3 marks]

Compare the coefficients for the three models.

Also visualise a few of the largest coefficients and the coefficients of features dropped by Lasso.

In [193...]

```
# Compare highest coefficients and coefficients of eliminated features
import pandas as pd

coeff_df = pd.DataFrame({
    "Feature": X_train_scaled.columns,
    "Coefficient": lasso_final.coef_
})

top_positive = coeff_df.sort_values(by="Coefficient", ascending=False).head()

top_negative = coeff_df.sort_values(by="Coefficient").head(10)

eliminated = coeff_df[coeff_df["Coefficient"] == 0]

print("===== TOP POSITIVE COEFFICIENTS (Strongest Price Boosters) =====")
display(top_positive)

print("===== TOP NEGATIVE COEFFICIENTS (Strongest Price Reducers) =====")
display(top_negative)

print("===== FEATURES ELIMINATED BY LASSO (Coefficient = 0) =====")
display(eliminated)
```

===== TOP POSITIVE COEFFICIENTS (Strongest Price Boosters) =====

	Feature	Coefficient
4	hp_kw	0.121955
35	make_model_Audi A3	0.054635
41	make_model_Renault Espace	0.052529
61	Gearing_Type_Semi-automatic	0.020237
34	make_model_Audi A1	0.020218
1	Gears	0.017365
17	Extras_other	0.010524
52	Fuel_Diesel	0.009052
49	Type_Nearly_New	0.008773
8	cons_comb	0.007948

===== TOP NEGATIVE COEFFICIENTS (Strongest Price Reducers) =====

	Feature	Coefficient
37	make_model_Opel Corsa	-0.132418
40	make_model_Renault Clio	-0.110660
2	age	-0.109603
0	km	-0.088445
60	Gearing_Type_Manual	-0.065041
36	make_model_Opel Astra	-0.059091
39	make_model_Other	-0.018404
57	Upholstery_type_Cloth	-0.015932
6	Displacement_cc	-0.010708
51	Fuel_Benzine	-0.008516

===== FEATURES ELIMINATED BY LASSO (Coefficient = 0) =====

	Feature	Coefficient
38	make_model_Opel Insignia	0.0
44	body_type_Sedans	-0.0
45	body_type_Station wagon	-0.0
53	Fuel_Other	0.0
54	Paint_Type_Metallic	0.0
59	Gearing_Type_Automatic	0.0
64	Drive_chain_front	-0.0

4 Conclusion & Key Takeaways [10 marks]

What did you notice by performing regularisation? Did the model performance improve? If not, then why? Did you find overfitting or not? Was the data sufficient? Is a linear model sufficient?

4.1 Conclude with outcomes and insights gained [10 marks]

📌 Conclusion & Key Takeaways — Regularisation Analysis

1. Did regularisation improve model performance?

Regularisation provided **minor improvements** in performance. Ridge Regression slightly improved RMSE and R² over Linear Regression, while Lasso reduced performance slightly but improved interpretability.

This shows the dataset already had a **strong linear signal**, so regularisation acted more as a stability enhancer than a performance booster.

2. Did regularisation reduce overfitting?

Yes — Ridge clearly reduced overfitting.

- Train and test R² scores were almost identical.
- Ridge performed consistently across alpha values.

This indicates the model generalises well and avoids overfitting.

3. Was there overfitting before regularisation?

Not significantly.

Linear Regression already produced:

- High R²
- Very small train–test performance gap

Regularisation simply refined the model rather than fixing major overfitting issues.

4. Was the dataset sufficient?

Yes, the dataset was sufficient and informative.

Evidence:

- Strong predictive performance (high R², low MAE/RMSE)
- Clear and monotonic relationships between predictors and car price
- Effective feature engineering improved model learning

The dataset supports linear modelling well.

5. Is a linear model sufficient for this problem?

Yes — a linear model is sufficient and effective.

Reasons:

- Residuals showed no major nonlinear patterns
- Features behaved linearly with respect to price
- Regularisation curves were smooth and stable
- No major heteroscedasticity or complexity requiring nonlinear models

Advanced models (Random Forest, XGBoost, Neural Nets) may offer small gains, but linear models already capture most of the variance.

6. What did regularisation actually achieve?

✓ Ridge Regression

- Stabilised coefficient values
- Improved generalisation slightly
- Was the best-performing model among the three
- Did not drastically change accuracy (model already strong)

✓ Lasso Regression

- Performed feature selection by eliminating weak predictors
 - Provided interpretability advantages
 - Slightly worse performance than Ridge but more concise model
-

★ Final Takeaway

The dataset was rich and well-structured, the linear assumptions held strongly, and the baseline model already performed extremely well. Regularisation improved **stability and interpretability**, not accuracy, because the model was already close to its optimal performance under a linear regime. A **linear model is fully adequate** for predicting used car prices in this scenario.

In []: